


COMS20001 - Concurrent Computing

www.cs.bris.ac.uk/Teaching/Resources/COMS20001

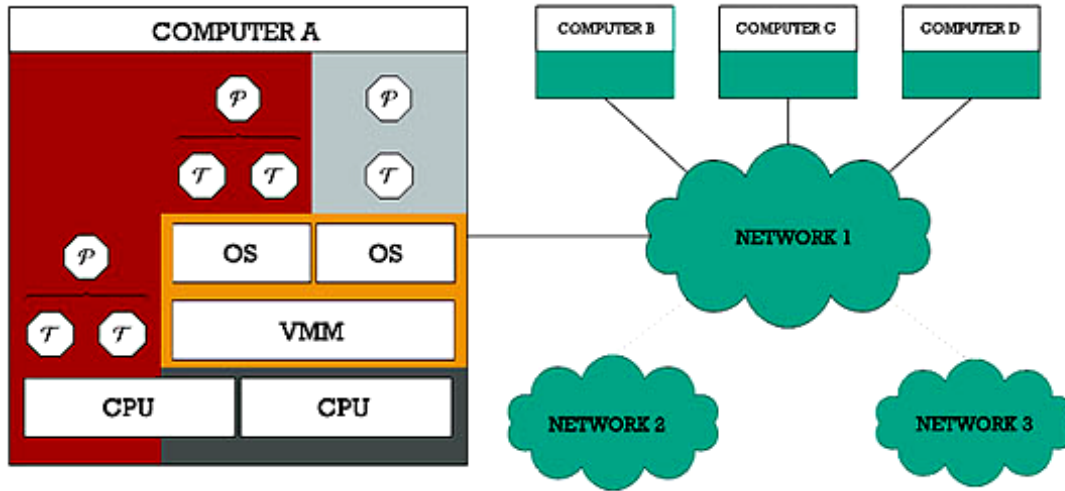
Lecture 01

Introduction and History

Sion Hannuna | sh1670@bristol.ac.uk
Dan Page | daniel.page@bristol.ac.uk



Unit Components of COMS20001



Year 1 Computer Architecture
Year 1 Programming & Algorithms

COMS20001 CONCURRENT COMPUTING:
PART 1: Concurrency Principles and Reasoning
PART 2: Concurrent Programming
PART 3: Operating Systems
PART 4: Interconnection Networks

Red	Weeks 01-12	Sion Hannuna
		Tilo Burghardt
Orange	13-18	Dan Page
Green	19-24	Dan Page




Concepts of Concurrent Computing
 Concurrent Programming in XC
 Formal Reasoning about Concurrency
 Operating Systems
 Interconnection Networks


What exactly is meant by **Concurrency**?

Concurrency is a property given to systems (e.g. software and/or hardware) in which **multiple sequential processes** are running **simultaneously, interacting** with each other.

SLOW



a traffic network



Brownian particle motion



the human brain

Examples of “naturally” concurrent systems

What are the implications of **Concurrency**?


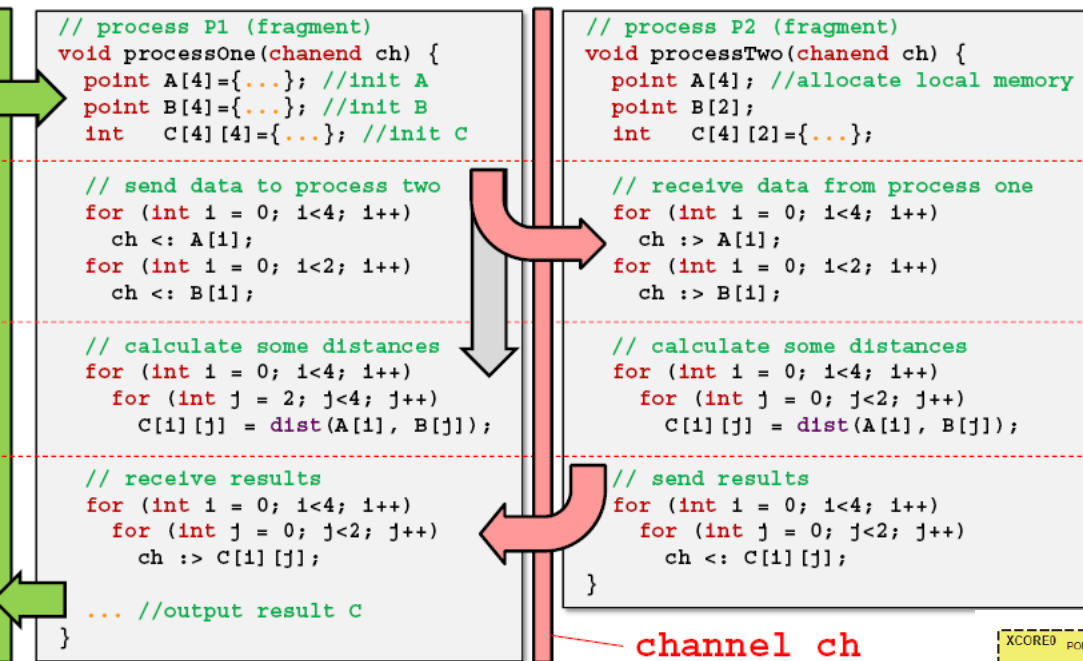
Theoretical/Conceptual aspects and models

- **History!**
- **Models of Process Communication:**
 - Shared memory, synchronous & asynchronous channels
 - Deadlocks
- **Paradigms of Hardware Parallelism**
 - Flynn's Taxonomy: SISD, SIMD, MISD, MIMD
- **Paradigms of Logical Parallelism**
 - Geometric Parallelisation
 - Parallelisation via Farming
 - Algorithmic Parallelisation
- **Performance issues**

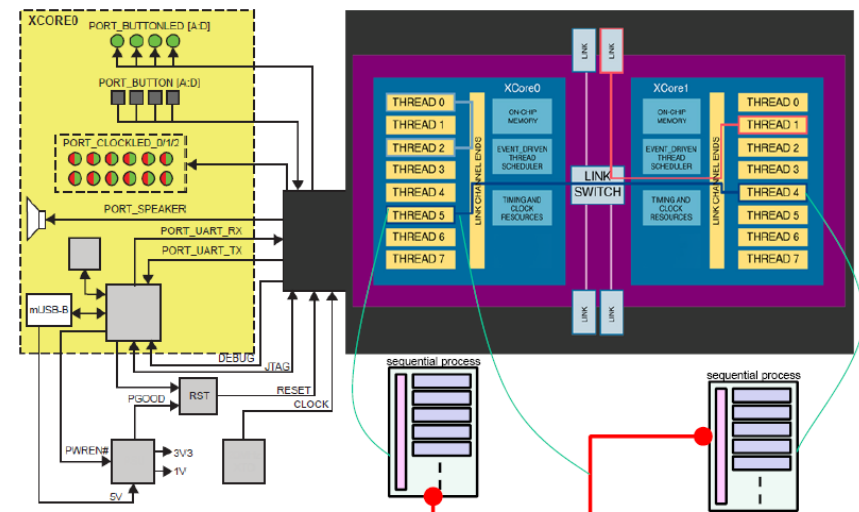
SLOW

Concurrent Programming in XC on xCore200

SION



- direct control of (parallel) hardware platform
- no operating system
- first hand experience of concurrency effects



Modelling with & Reasoning about Concurrent Systems

SION
& TILO

Failures of a process:

$failures(P) = \{(tr, X) \mid tr \in traces(P) \text{ and } X \in refusals(P/tr)\}$

■ $P = a \rightarrow b \rightarrow STOP$ with $\alpha(P) = \{a, b\}$

Transition Diagram of P :



$traces(P) = \{ \langle \rangle, \langle a \rangle, \langle a, b \rangle \}$

$refusals(P/\langle \rangle) = \{ \{\}, \{b\} \}$

$refusals(P/\langle a \rangle) = \{ \{\}, \{a\} \}$

$refusals(P/\langle a, b \rangle) = \{ \{\}, \{a\}, \{b\}, \{a, b\} \}$

$failures(P) = \{ (\langle \rangle, \{\}), (\langle \rangle, \{b\}), (\langle a \rangle, \{\}), (\langle a \rangle, \{a\}), (\langle a, b \rangle, \{\}), (\langle a, b \rangle, \{a\}), (\langle a, b \rangle, \{b\}), (\langle a, b \rangle, \{a, b\}) \}$

Conditions for Deadlock

Coffman, Elphick and Shoshani identified 4 *necessary and sufficient* conditions for deadlock [System Deadlocks. ACM Computing Surveys 3, 2 (June), p. 67-78, 1971.]

1. Agents claim exclusive control of the resources they require.

⇒ “Mutual exclusion”

2. Agents hold resources that prevent other agents from obtaining additional resources.

⇒ “Wait for” condition

3. Resources cannot be released by the agents holding them until the resources are held by other agents.

⇒ “No preemption” condition

4. A circular chain of agents exists, each of which is waiting for resources that are held by the next agent in the chain.

⇒ “Circular wait” condition

Definition of Petri Nets

Formally, a Petri Net N is a 4-tuple describing an annotated, directed, bipartite graph:


$$N = \{P, T, A, M_0\}$$

where

- P is a finite set of places
- T is a finite set of transitions
- A is a finite set of arcs (arrows)
- M_0 is the initial token marking



Carl Adam Petri



Elements of a Petri net

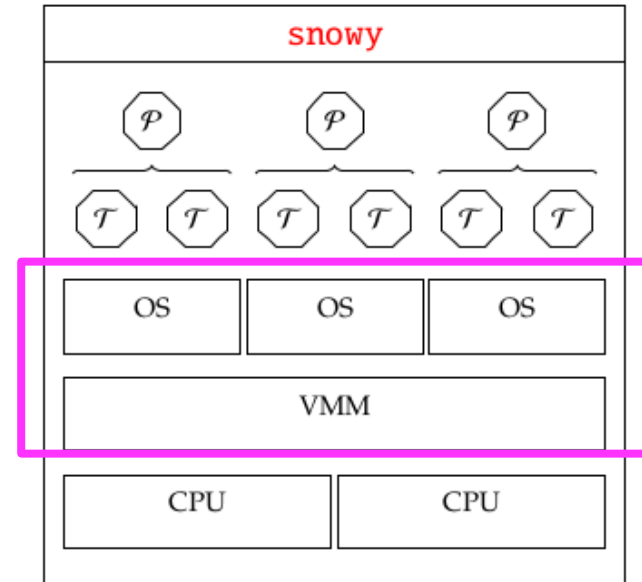
- formal modelling of concurrent processes using CSP basics
- reasoning about communicating sequential processes (e.g. safety, liveness, deadlock, divergence)
- understanding the impact of design choices on program properties (e.g. fairness, starvation, overheads)

Operating System Concurrency


The OS is the “Middle-man” of a computing system. It coordinates concurrent operations on a hardware platform.

There are many types of concurrent operations which are managed by an operating system, including:

- Process management
- Memory management
- Device management



Dan



Operating System Concurrency

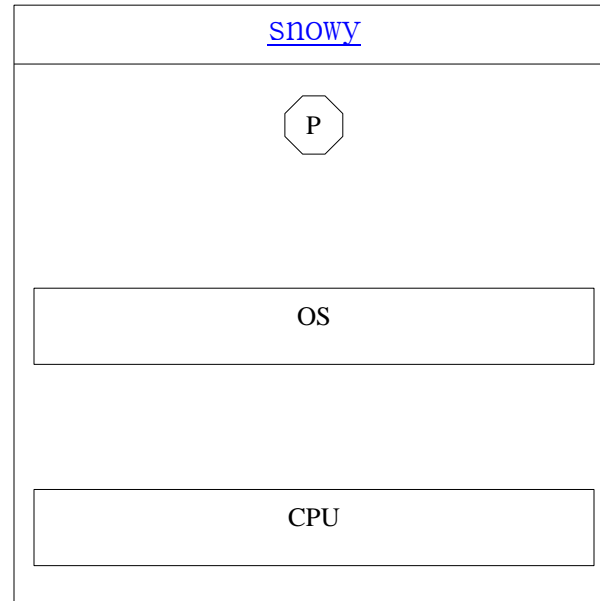
Dan

In the OS part of the course, we will:

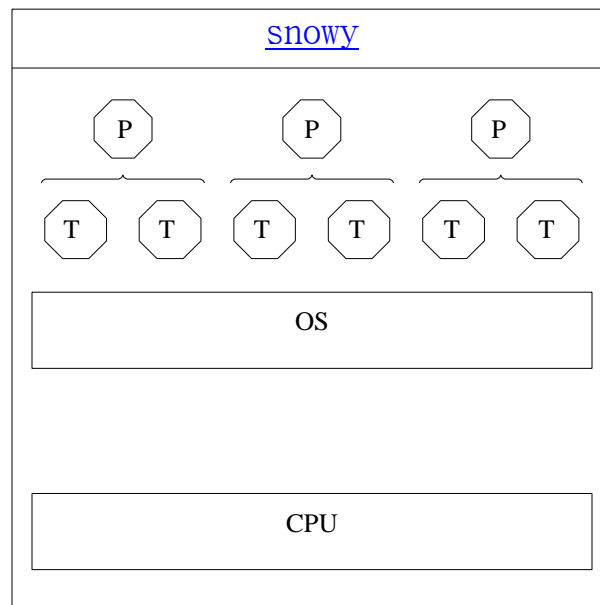
- Look at what concurrency resources are provided by hardware
- Investigate how these are managed by an OS to provide concurrent execution
- Build a multi-tasking OS



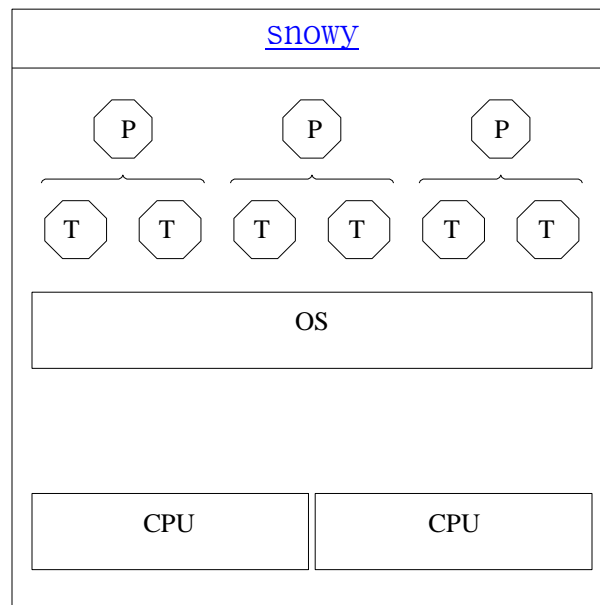
Concept of concurrent computing exists at various scales



Concept of concurrent computing exists at various scales

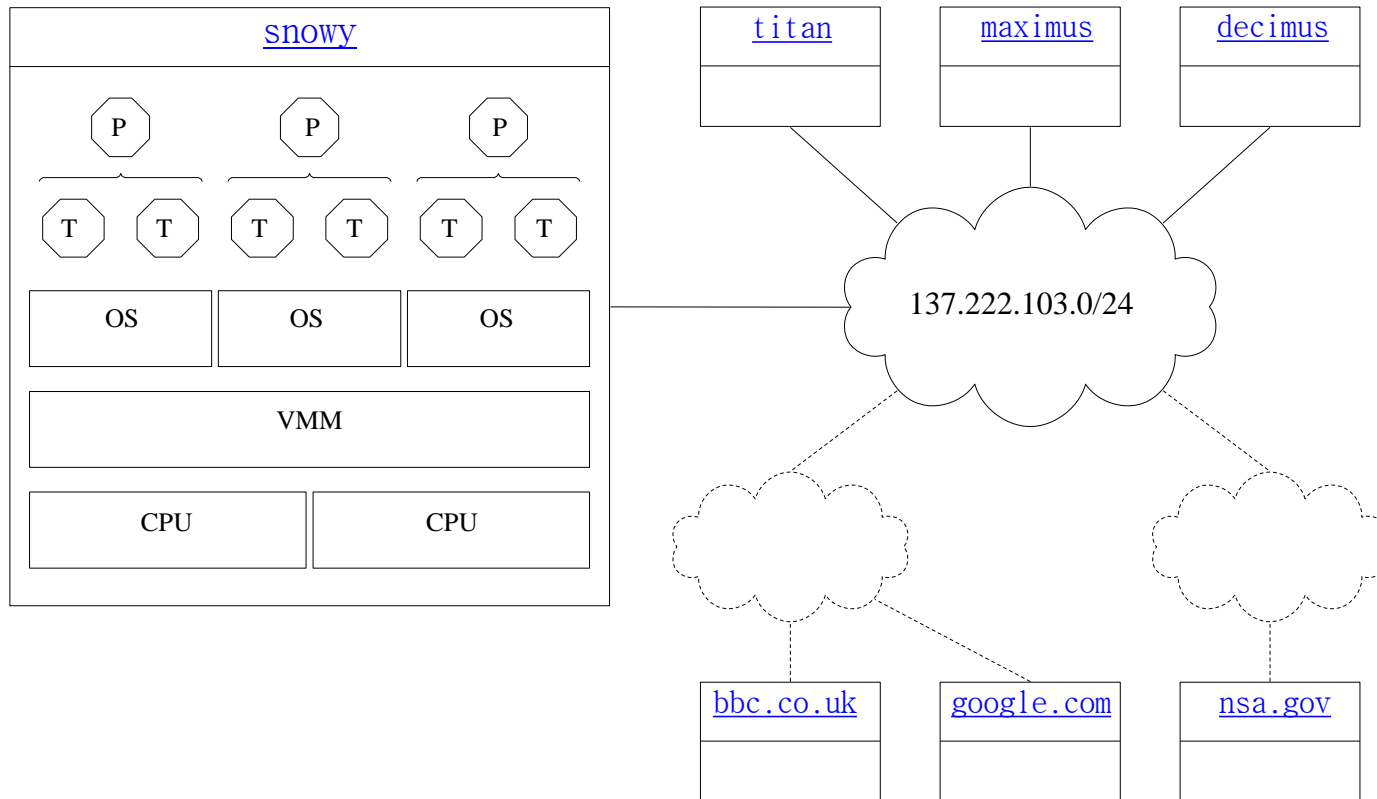


Concept of concurrent computing exists at various scales



Concept of concurrent computing exists at various scales

Dan



Bottom-up Understanding of Network Environments

Dan

1. link layer
transmission, modulation, multiplexing, metrics
802.3 (Ethernet), 802.11 (WiFi)
2. internet layer
addressing, forwarding, fragmentation
IP, ICMP, DHCP, ARP
3. transport layer
connection management, routing, flow and congestion control
UDP, TCP, NAT, DNS
4. application layer

... while emphasising **general** concepts as applied in **specific** technologies (using a running example) ...

... or, explain how the Internet works (and doesn't) so you can

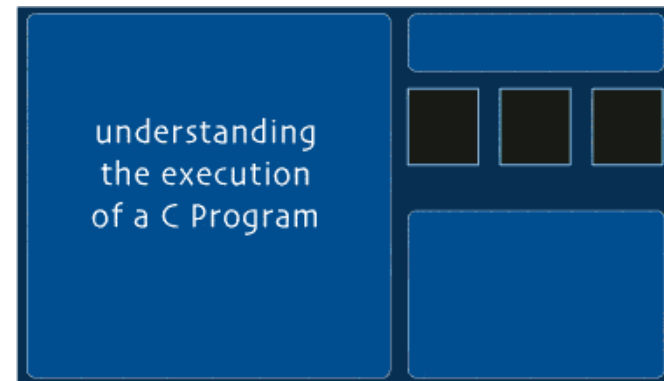
1. make effective use of existing technologies, but also
2. understand and apply similar concepts in next-generation technologies (of your own design).

So far on your degree programme...

- ...you learned about principles and programming basics...
 - C (procedural programming),
 - Java (object oriented programming) and
 - Haskell (functional programming)
- ...however, you essentially wrote **sequential** programs...
 - where programs were **deterministic** entities (single, defined sequence of state changes)
 - where **memory** (code & data) was accessed **sequentially** (instruction by instruction) and **available globally** (apart from scope)


IN THIS PART OF COMS20001...

What happens if multiple programs are run simultaneously on multiple CPUs, with or without shared resources (memory banks etc)?

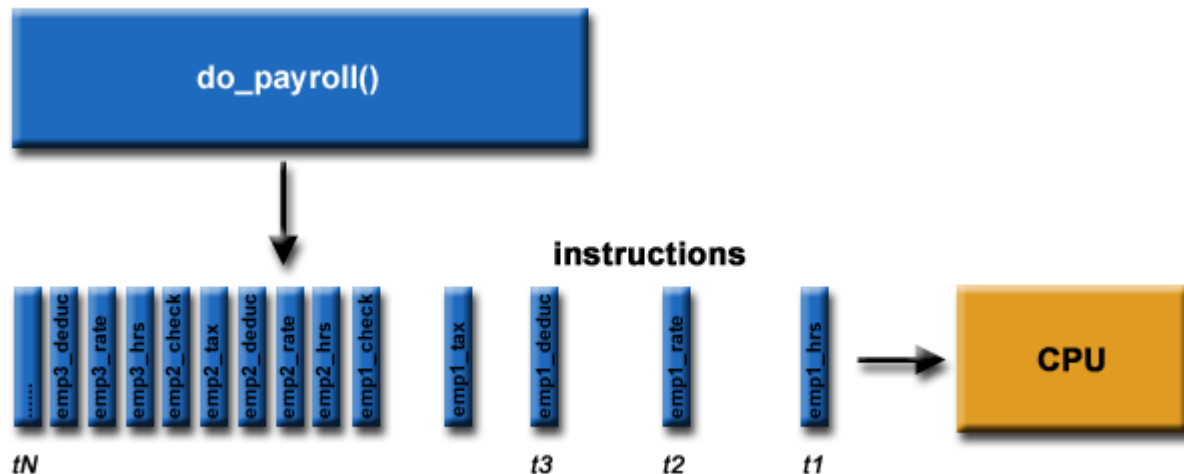


©2004 HowStuffWorks

So essentially, we want to go from this...




For example:



To this...


For example:



Disclaimer: this is a simple scenario – in practice things are a lot more complex...


What is a 'Sequential Process'?

- A ***sequential process*** is a unit of sequential execution.
- We structure complex systems as sets of simpler activities, each represented as a ***sequential process***.
- So processes are the basic building block of concurrent systems.
- Processes can overlap or be parallel, so as to reflect the concurrency inherent in the physical world, to offload time-consuming tasks, to manage communications or other devices




Is it really parallel?

3 parallel processes running . . .



. . . on the same CPU



. . . on 3 different CPUs

Concurrency? Parallelism?

- Concurrency
 - Logically simultaneous processing.
 - Does not imply multiple processing elements (PEs).
 - Interleaved execution on a single PE.
- Parallelism
 - Physically simultaneous processing.
 - Involves multiple PEs
 - Independent device operations.

Both concurrency & parallelism require controlled access to shared resources.

We shall use the terms parallel and concurrent interchangeably, and will only distinguish between real and pseudo-parallel execution if necessary.

Why is 'Concurrency' essential today?

Multiprocessing hardware → parallelism

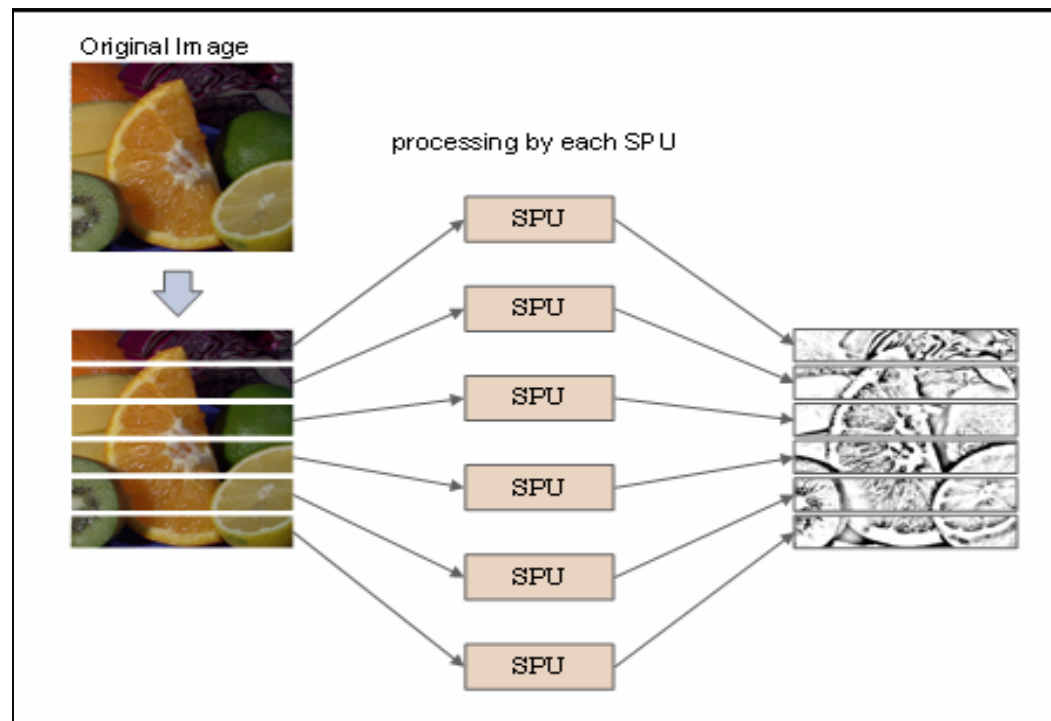
- Performance gain
- Increased application throughput
 - an I/O call need only block one thread
- Increased application responsiveness
 - e.g. high priority processes for user requests...
- More appropriate structure for multitasking applications
 - e.g. one which interacts with the environment, controls multiple activities and handles multiple events.

Concurrency is becoming hard to ignore!

- Multicore chips are increasingly prevalent
 - Only multithreaded applications will see the performance benefits these chips offer
- Programming for the Web often requires concurrent programming (think image loading in Web browsers)
- Lots of domains in which concurrency is the norm, e.g.
 - Embedded software systems
 - Robotics
 - Simulation and modelling (e.g. weather prediction systems)

Example: image processing

- Trade *processing over time* for *processing over space*
- Performance gain from multiprocessing hardware



Must handle concurrency properly!

Write efficient concurrent programs

- ...understand **fundamental causes** why programs can produce different behaviours/results when run concurrently (e.g. deadlocks, race conditions, synchronisation etc)
- ...consider **systematic approaches** to control effects of concurrency: avoid 'freezes' and 'unresponsiveness' and optimize parallel execution performance


Therac-25: Computerized radiation therapy machine's concurrent programming errors contributed to accidents causing deaths and serious injuries

Mars Rover: Problems with interaction between concurrent tasks caused periodic software resets reducing availability for exploration

Brief History of Concurrent Computing

It all started so innocently...

- until 1950s: Single-User Mode & Batch Processing
 - Single-User Mode: reserving the entire machine
 - e.g. 'I book from 9:00-13:00...'
 - debugging time was included in reserved time slot
 - Batch Processing: submitting jobs to a 'queue'
 - when program ends/fails next job is immediately executed
 - CPU still idle during I/O , low responsiveness of system



Colossus Mark 2 (1940s)



Universal Automatic Computer UNIVAC (1951)

First Steps towards Simultaneous Execution...

- 1960s: Era of Multi-Programming
 - multiple programs reside in memory at once
 - switching between programs using interrupts
 - introduction of concepts 'memory protection' and 'virtual memory'
 - beginning of 'theory of concurrent systems'
 - How to synchronize the processing of all these programs smoothly?
 - Communicating Sequential Processes (CSP - discussed later in the unit)



Leo I mainframe computer (1960s)



IBM mainframe (late 1960s)

Rise of Resource-Sharing Operating Systems

- 1970s: Time-Sharing and Memory-Sharing via OS
 - Fast Context Switching: time-slicing of a CPU between programs
 - illusion of using a (perhaps slower) machine alone
 - options of prioritising programs (in terms of resource use, rights etc)
 - but: switching overhead is significant
 - Shared Memory: programs share memory to communicate
 - enables information transfer between programs
 - but: mixing concepts of data storage and program communication



Mission control, Houston (1970s)




Apple II (1977)


Hardware Performance drives Efficiency Gains

- 1980s: Era of Superscalar expansion (instruction level parallelism)
 - 50% annual improvement in performance
 - pipeline processor (10 CPI \rightarrow 1 CPI) for implicit parallelism
- 1990s: Era of Diminishing Returns
 - branch prediction etc. (1 CPI \rightarrow 0.5 CPI)
 - performance below expectations


cycles per instruction

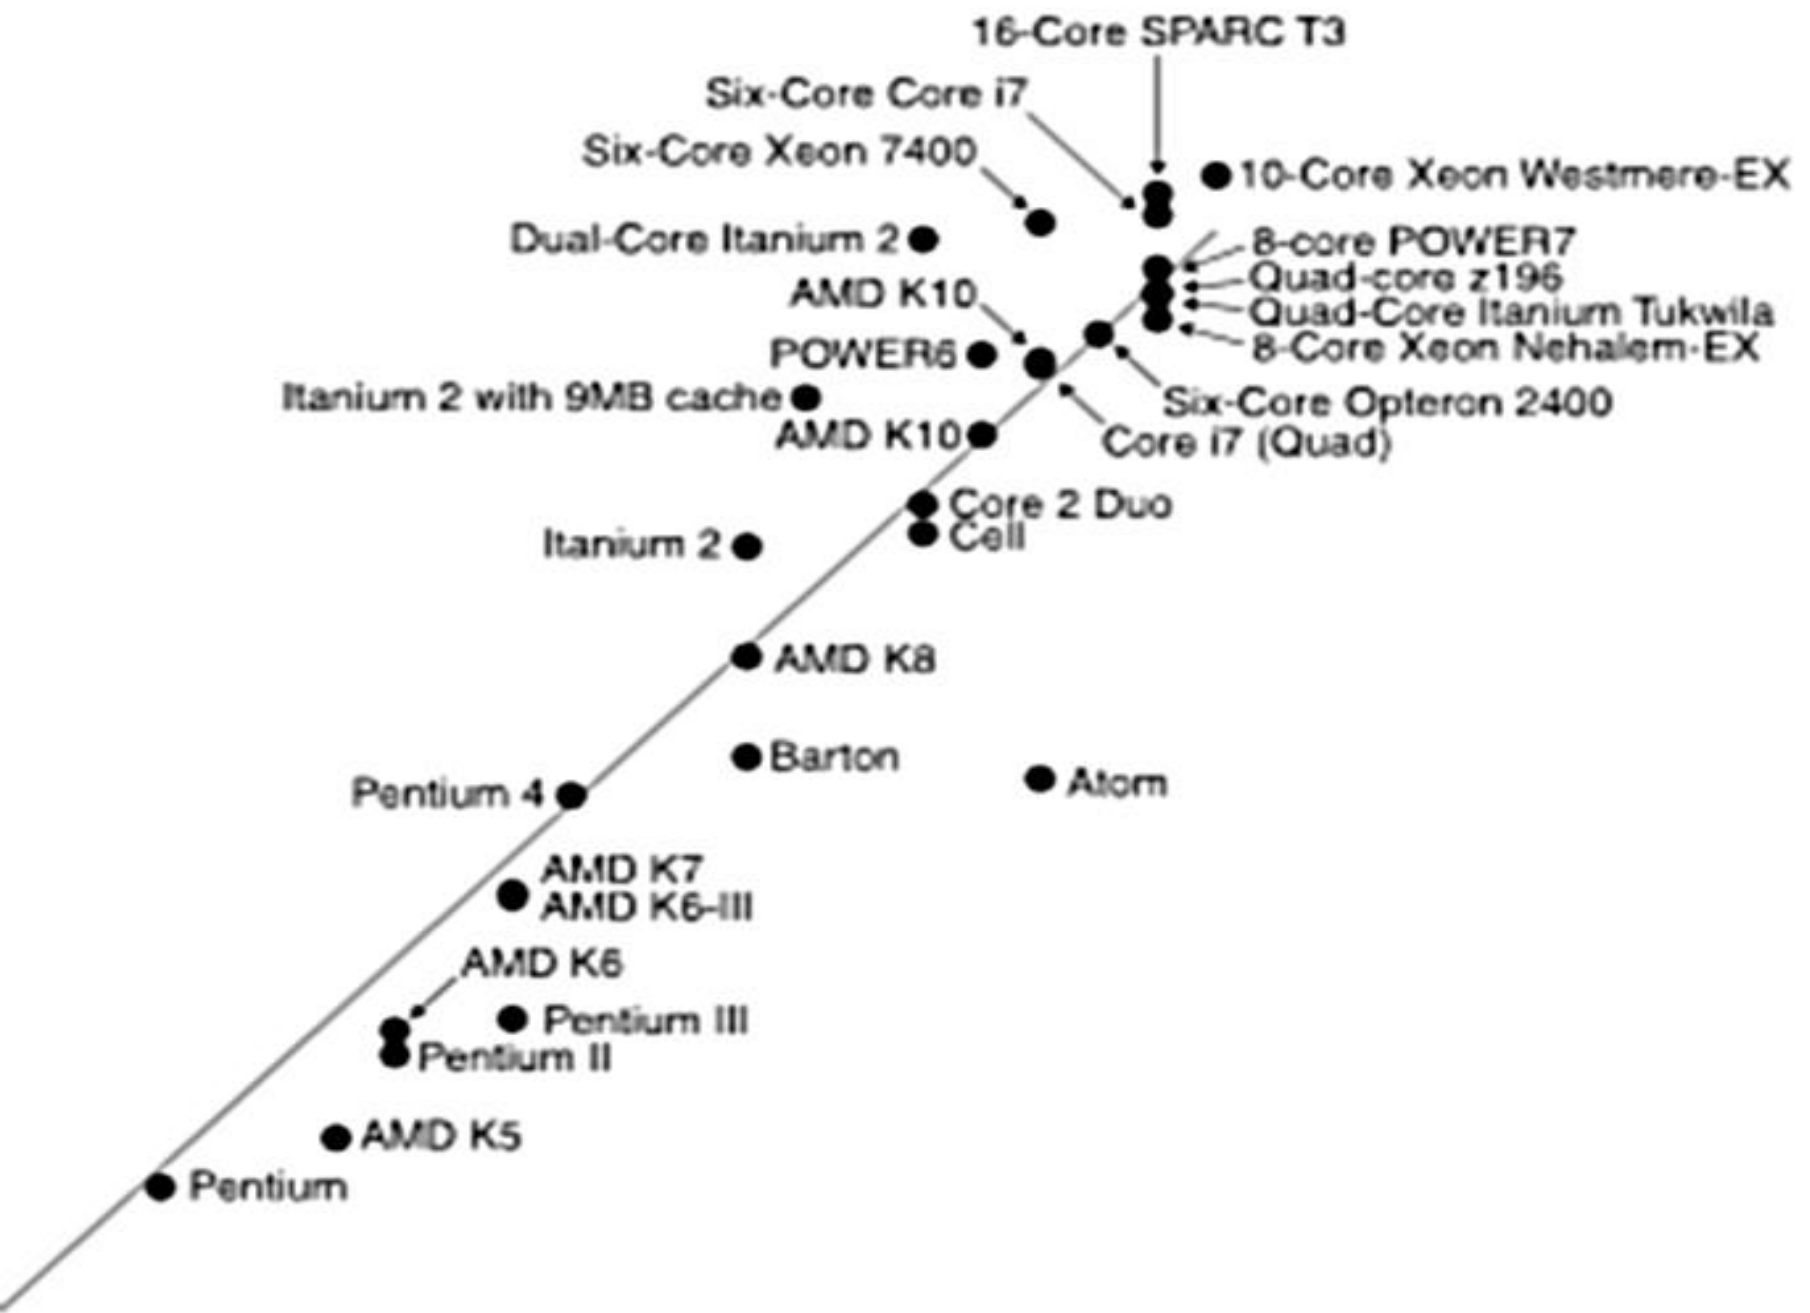


CRISIS: Clock Speeds hit Ceiling in 2000s




Microprocessor Transistor Counts 1997-2011 & Moore's Law





A Multi-Core Revolution

Concurrent Programming is Back on the Menu!



Therefore: Concurrency is critical now!

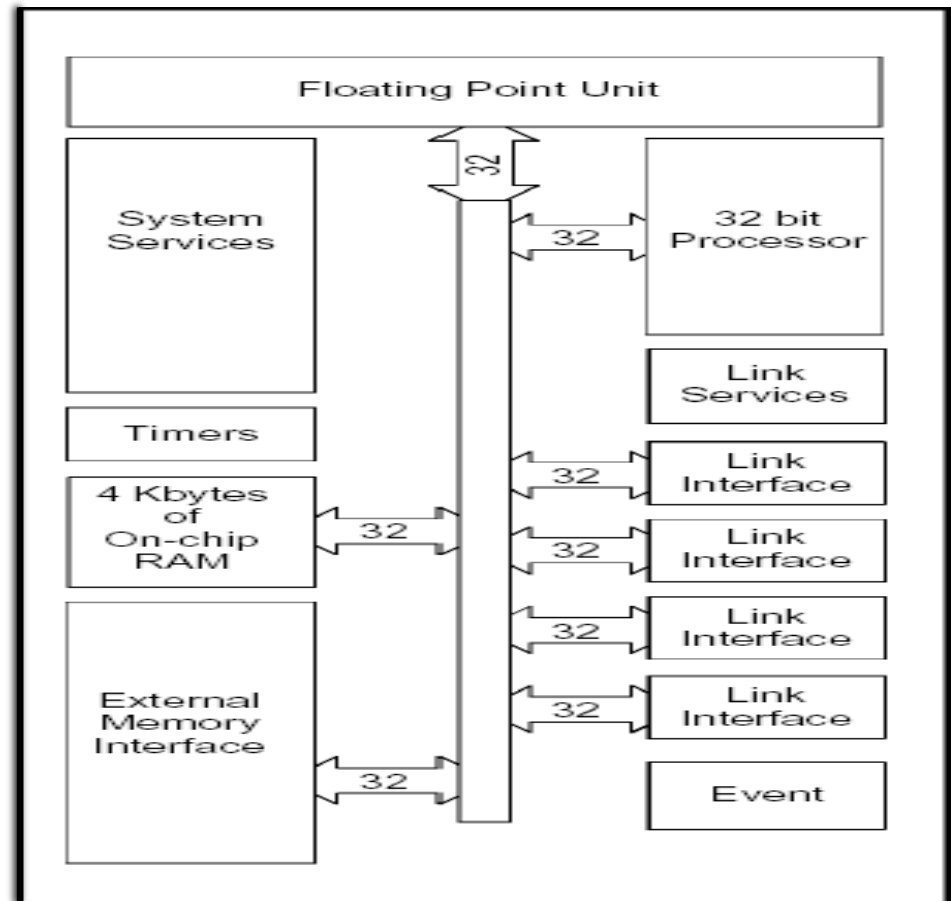

Multiprocessing hardware → parallelism

- Performance gain
- Increased application throughput
 - an I/O call need only block one thread
- Increased application responsiveness
 - e.g. high priority processes for user requests...
- More appropriate structure for multitasking applications
 - e.g. one which interacts with the environment, controls multiple activities and handles multiple events.


Old News: It has been there a long time

The Transputer – 1980s

- The first single chip computer designed for message-passing parallel systems, in 1980s, by INMOS
- *Transistor Computer*. Low cost, low power chip to form a complete processor
- It had a RISC type of instruction set




Transputer Network




Transputer Network Board

Transputer networks with specific connections



Technology constantly on the move!

- The Intel® Core™ i7 microprocessor (“Nehalem”)
 - 4 cores/chip
 - 45 nm, Hafnium hi-k dielectric
 - 731M Transistors
 - Shared L3 Cache - 8MB
 - L2 Cache - 1MB (256K x 4)
- Num of transistors not limiting factor
 - Currently ~ 2.5 billion transistors/chip
 - Problems:
 - Too much Power, Heat, Latency
 - Not enough Parallelism
- 3-dimensional chip technology?
 - Sandwiches of silicon
 - communication can be an issue



There is a Software side to the story...

First Software Crisis (1960s-1970s)

Problem: assembly language limits abstraction and portability

Solutions: **high-level languages** for uni-processors (C, Fortran etc)

Second Software Crisis (1980s-2000)

Problem: procedural languages limit maintainability and reusability

Solutions: **object orientation** (C++, Java, C# etc), software design process and tools

NOW: Third Software Crisis (2005-today)

Problem: sequential approach limits data throughput, real-time demand and distributed infrastructure

Challenge: explicit support for concurrency and communication...

- high-performance connectivity (Gbit networks, 4G, ...)
- use physical parallelism (multicore, multiCPU, cloud...)
- **logical parallelism** (concurrent programming...)

Sharing Memory vs. Process Communication

Traditional: *Shared Memory Model*

- concurrent interaction via synchronized alterations of shared memory (e.g. C#, Java)
- problem of synchronization (controlling the concurrency) is left to the responsibility of the programmer!



Makes writing concurrent programs much more difficult than writing sequential programs!

More common Approach: *Message Passing Model*

- processes communicate by exchanging messages on channels (e.g. Occam, Ada, XC)
- can be efficiently implemented in existing multi-processor hardware with *or without* shared memory

Languages based on Parallel Execution: Occam

Occam: 1st language based on principles of **parallel execution**.

- automatic communication and synchronization between concurrent processes.
- Theoretical foundation based on CSP, first introduced by Tony Hoare in 1968. (CSP = Communicating Sequential Processes)
- Inmos developed Occam and the Transputer in the 1980s.

Philosophy behind the Occam language:
Minimalistic approach. Easy to learn.

14th century Franciscan friar William of Ockham:
“It is vain to do with more what can be done with fewer.”



Languages based on Parallel Execution: XC

- C-based imperative language developed by XMOS
- shows semantic similarities to Occam in a familiar C syntax
- design was heavily influenced by CSP, which can be used to reason about XC programs
- supports explicit parallelism and channel communication (for versatility: it also includes guarded commands)
- directly supports timers and port access
- compiles directly to drive XMOS multi-core hardware

You will get your own xmos board kit for your coursework later



TASK: Form Pair Programming Teams Now

- Within the next 48 hours: find a team partner and register on the course website for one of the two lab slots
- You will then work together during TB1 **as a pair** in your programming team



Lecture 2 Outlook



XC: Towards Concurrent Programming in xC