

COMS20001 - Concurrent Computing

www.cs.bris.ac.uk/Teaching/Resources/COMS20001



Lecture 05

Replication & Pipelining

Sion Hannuna | sh1670@bristol.ac.uk
Dan Page | daniel.page@bristol.ac.uk

Process Replication using Replicated PAR

- builds an array of **structurally similar** parallel processes
- parameterise processes using the replicator index (e.g. *i*)
- **Example:** use together with channels and arrays for process chains, buffers, queues etc.

```
// chain of processes
```

replication.xc


```
...  
void chainElement( chanend cInput, chanend cOutput) { ... }  
  
int main ( void ) {  
    chan c[4];  
    par (int i=0; i<4; i++)  
        chainElement(c[i], c[(i+1)%4]);  
    return 0;  
}
```

Replicated **PAR** is key to elegant concurrent programming!

Example: Process Buffer via Replicated PAR

```
...  
chan queue[11];           // array of 11 channels  
par (int i=0; i<10; i++) { // 10 parallel processes p in buffer  
    int a;  
    queue[i] := a;  
    p(i, a);  
    queue[i+1] <: a;  
}  
...
```

Is this legal (it's not)?



NOTE!

Still need processes to feed the queue and to bleed the queue

```

...
void inputProcess(chanend c){ //FEEDING buffer pipeline
    while (1)
        c <: 1; //send input data item, e.g. key from keyboard, sensor data etc...
}

...
void outputProcess(chanend c){ //DRAINING buffer pipeline
    int x;
    while (1) {
        c :=> x; //drain element from last buffer element
        printf("%d",x); //display on screen
    }
}

...
void bufferProcess(chanend cIn, chanend cOut){ //BUFFER one element
    int x;
    while (1) {
        cIn :=> x; //try to read (and hold, i.e. buffer)
        cOut <: x; //try send data item on to next buffer
    }
}

...
int main(void) { //SETUP CONCURRENT PROGRAM
    chan queue[14]; //create 14 channels
    par { //align 15 concurrently running processes in pipeline
        on tile[0]: inputProcess(queue[0]); //start first process of pipeline
        on tile[1]: outputProcess(queue[13]); //start last process of pipeline


        par (int i=0;i<13;i++) { //replicate 13 buffer processes
            on tile[i%2]: bufferProcess(queue[i],queue[i+1]);
        }
    }
    return 0;
}
...

```

Example:
Buffer via
Replication
and Nesting
with **PAR**

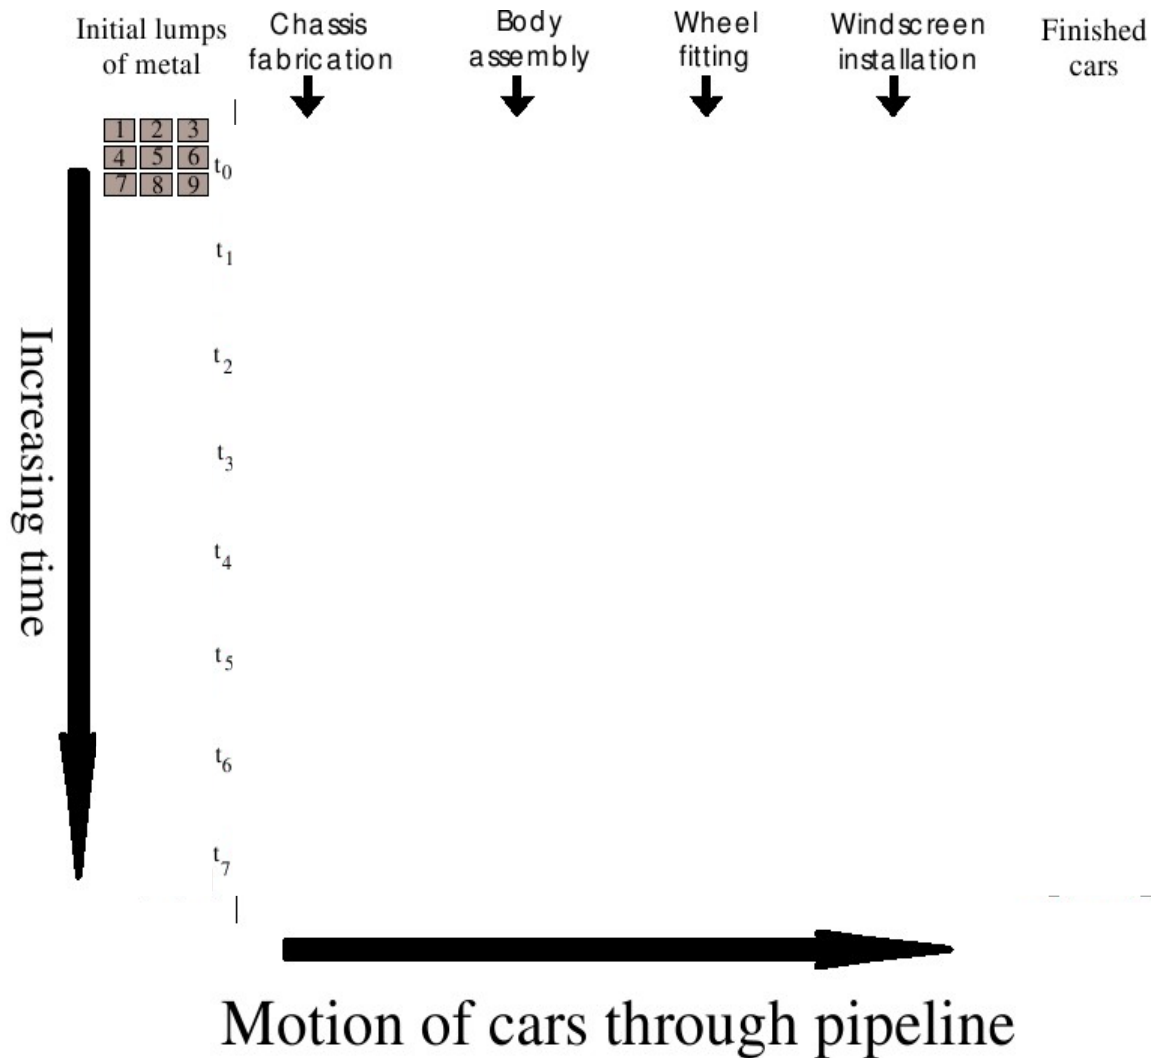
Pipeline Processing

- Pipeline stages transform the data before passing it on in a chain. Can be (much) faster than using a loop.
- Speed advantage of pipeline is due to overlapping in time



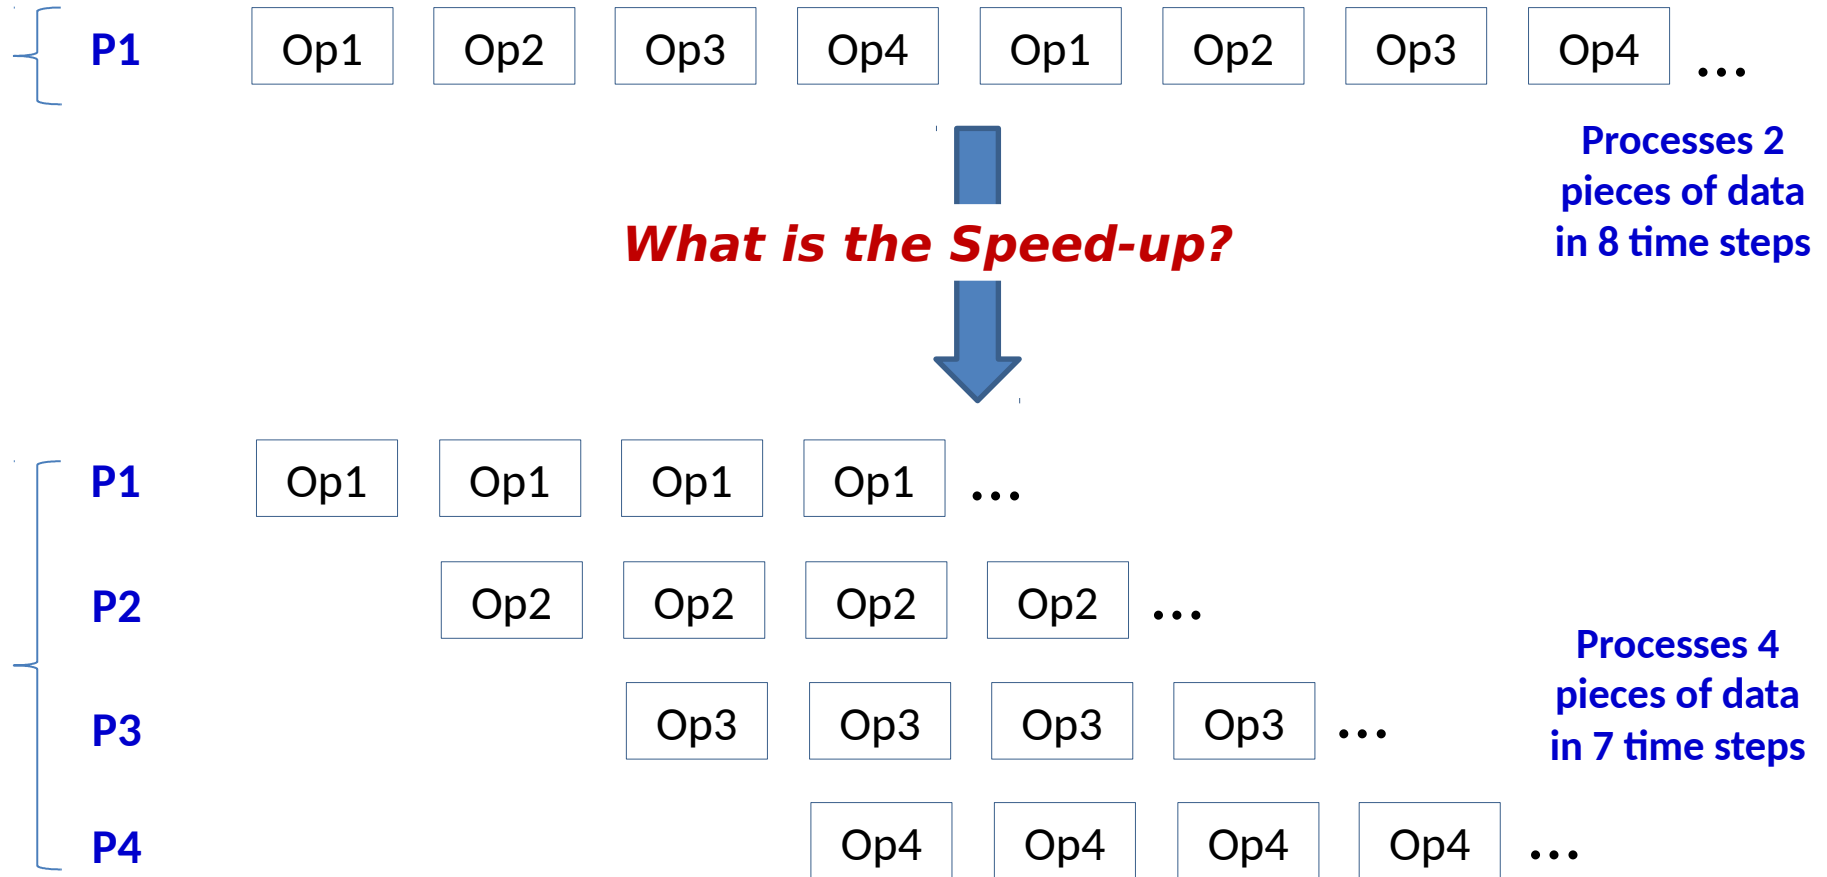
Max time-saving achieved if each stage runs on its own processor.

Pipelines: Car making analogy



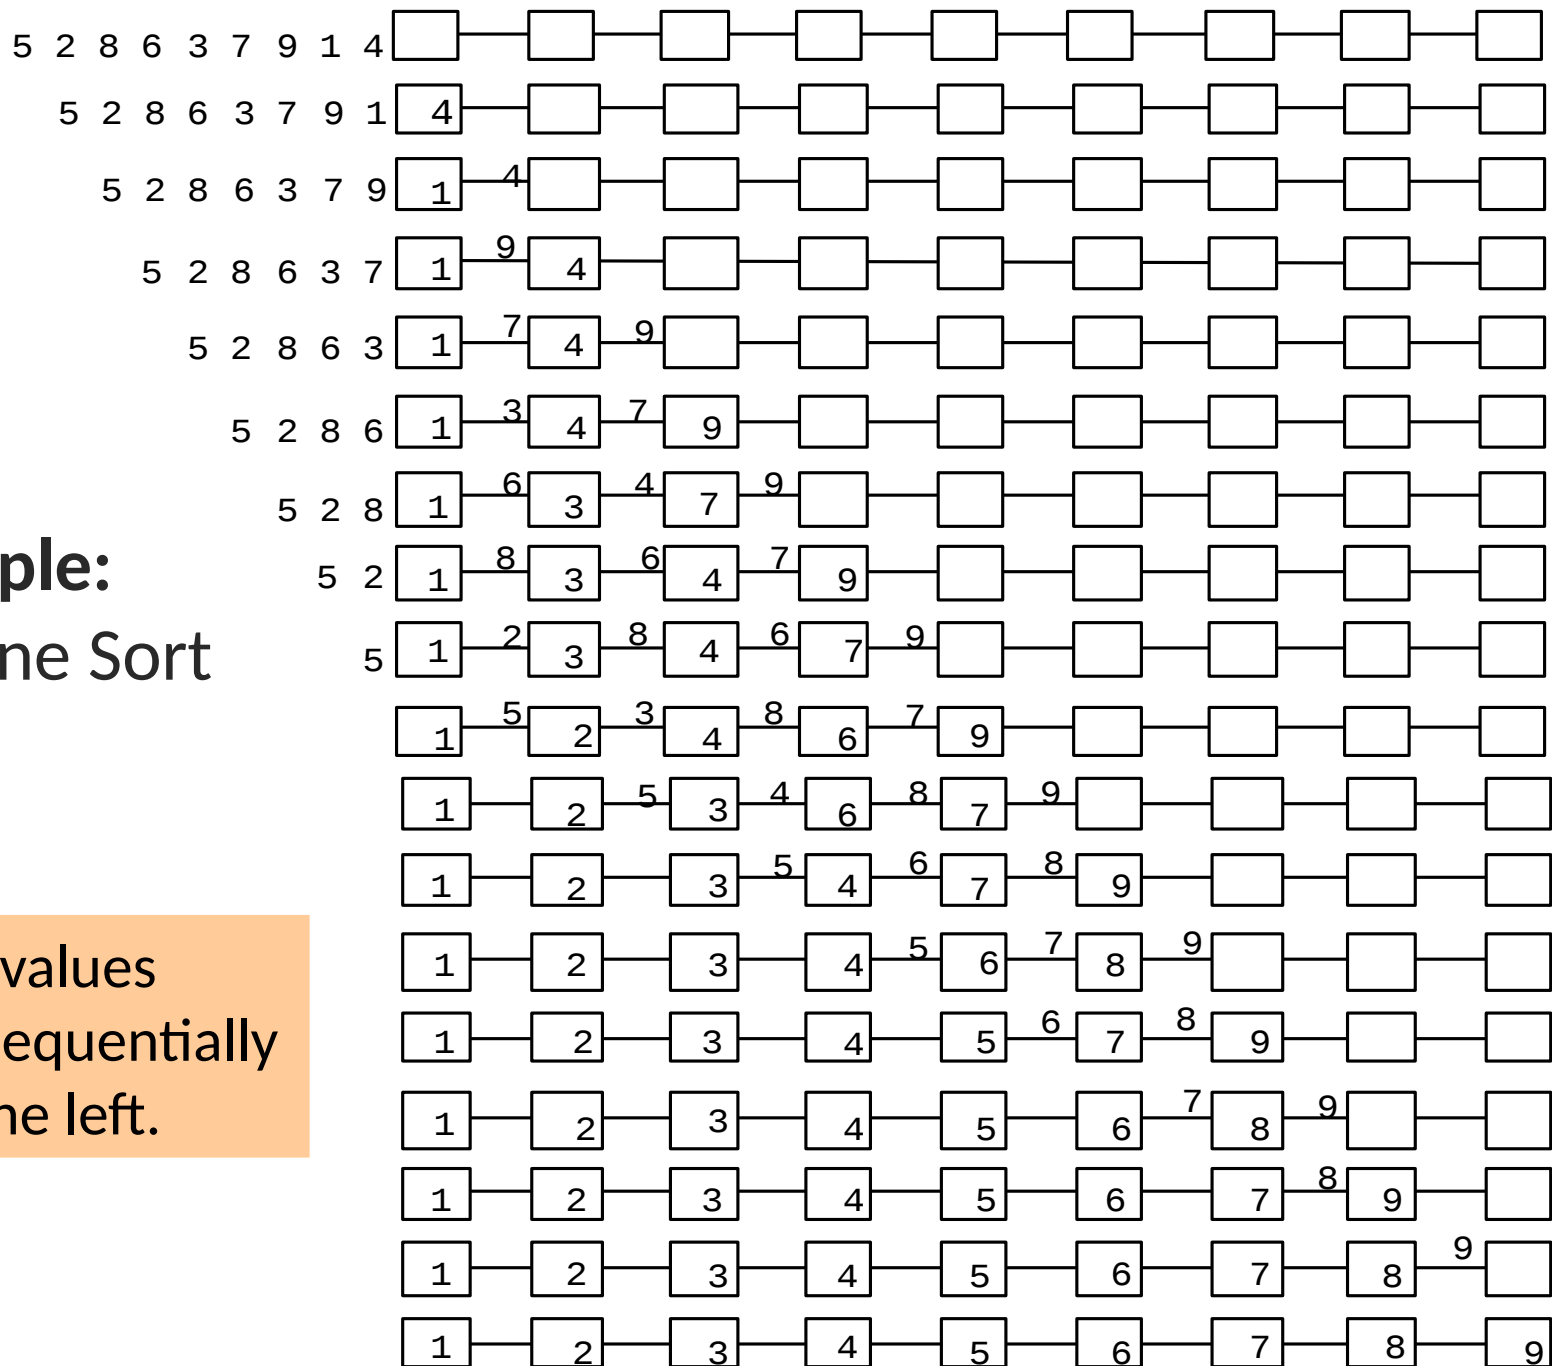

From "Alan Chalmers, Practical Parallel Processing, 1996"

Pipeline Processing: speed-up?



We needed 4 processors to get only 2x speed-up.

More on this later...



Example: Pipeline Sort

...with values
input sequentially
from the left.

Example: Pipeline Sort

Sort a stream of **N** numbers with a pipeline of **N** processes!

- Each process has two local variables: **lowest** and **next**.
 - Each number enters a process and is compared to value in **lowest**.
 - If not smaller than **lowest** then pass to next pipeline stage
 - If smaller than **lowest** then keep it and pass on previous lowest
 - When all numbers are processed, pass final value of **lowest** on
- ⇒ Develop code using Replicated **PAR**!

XC Pipeline Sort Code I

- To sort a sequence of 10 numbers:

```
// XC code to compare numbers for each pipeline stage
...
int next;
cIn :> next;
if (next >= lowest) {
    cOut <: next;
}
else {
    cOut <: lowest;
    lowest = next;
}
...
```

XC Pipeline Sort Code II

- Now repeat comparison code 9 times in each pipeline stage:

```
// XC code for one single pipeline sort stage
```

```
void sortStage(chanend cIn, chanend cOut) {  
    int lowest;  
    cIn :=> lowest;  
    for (int j=0;j<9;j++) {  
        int next;  
        cIn :=> next;  
        if (next >= lowest) {  
            cOut <:= next;  
        }  
        else {  
            cOut <:= lowest;  
            lowest = next;  
        }  
    }  
    cOut <:= lowest;  
}
```

XC Pipeline Sort Code III

```
// XC code fragment for pipeline sort chain
...
void sortStage(chanend cIn, chanend cOut) {
    int lowest;
    cIn :> lowest;
    for (int j=0;j<9;j++) {
        int next;
        cIn :> next;
        if (next >= lowest) {
            cOut <: next;
        }
        else {
            cOut <: lowest;
            lowest = next;
        }
    }
    cOut <: lowest;
}
...
int main(void) {
    chan pipe[11]; ...
    par (int i=0; i<10; i++) {
        on tile[i%2]: sortStage(pipe[i], pipe[i+1]);
    }
    return 0;
}
```

- Now replicate 10 times to get 10 parallel processes, one for each pipeline stage
- Note, we need one channel more than the number of processes:
chan pipe[11];

Main
program

XC Pipeline Sort Code IV

- Processes to feed and bleed the pipeline

```
// XC code to feed and bleed pipeline
```

```
...
```

```
void inputProcess(chanend c) {  
    for (int j=0;j<10;j++)  
        c <: readUnsortedNumber(j);  
}
```

```
...
```

```
void outputProcess(chanend c) {  
    for (int j=0;j<10;j++) {  
        int x;  
        c :> x;  
        printf("%d, ", x);  
    }  
}
```

```
...
```

Full XC Pipeline Sort Code

```

void inputProcess(chanend c) { //FEEDING pipeline
    for (int j=0;j<10;j++)
        c <: readUnsortedNumber(j);
}

...

void outputProcess(chanend c) { //BLEEDING pipeline
    for (int j=0;j<10;j++) {
        int x;
        c :=> x;
        printf("%d",x);
    }
}

...

void sortStage(chanend cIn, chanend cOut) { //one STAGE of pipeline
    int lowest;
    cIn :=> lowest;
    for (int j=0;j<9;j++) {
        int next;
        cIn :=> next;
        if (next >= lowest) {
            cOut <: next;
        }
        else {
            cOut <: lowest;
            lowest = next;
        }
    }
    cOut <: lowest;
}

...

int main(void) { //SETUP CONCURRENT PROGRAM
    chan pipe[11]; ...
    par {
        inputProcess(pipe[0]);
        outputProcess(pipe[10]);
        par (int i=0; i<10; i++) {
            on tile[i%2]: sortStage(pipe[i],pipe[i+1]);
        }
    }
    return 0;
}

```

More Efficient XC Pipeline Sort Code

pipelinesort2.xc

```
...  
void sortStage(int i, chanend cIn, chanend cOut) {  
    int lowest;  
    cIn :> lowest;  
    for (int j=0;j<9-i;j++) { //sort unsorted part  
        int next;  
        cIn :> next;  
        if (next >= lowest) {  
            cOut <: next;  
        }  
        else {  
            cOut <: lowest;  
            lowest = next;  
        }  
    }  
    cOut <: lowest;  
    for (int j=0;j<i;j++) { //copy already sorted part  
        cIn :> lowest;  
        cOut <: lowest;  
    }  
}  
...  
int main(void) {  
    chan pipe[11];...  
    par (int i=0; i<10; i++) {  
        on tile[i%2]: sortStage(i,pipe[i],pipe[i+1]);  
    }  
    return 0;  
}
```

TAKE HOME
EXERCISE

- Data already sorted by previous pipeline nodes
- ▢ do not sort in these cases, but simply copy data

Parallelism suitable for Pipelining

In general, a program that would otherwise be written using two or more nested **WHILE** loops can often be transformed into a pipelined program.

```
WHILE test1
  WHILE test2
    ...
```

No speed advantage unless the inner loop produces a value to pass on early in its execution.




if value is only produced at termination, then there is no overlap,
hence no parallelism then to be exploited in a pipeline.

Simple Puzzle: Traditional Pipeline Concept


Dirty Laundry

Ann, Brian, Cathy, Dave
each have one load of clothes
to wash, dry, and fold


- Washer takes 30 minutes
- Dryer takes 40 minutes
- “Folder” takes 20 minutes



Sequential Approach




Simple Pipeline Approach



If they learned pipelining, how long would laundry take and so what time would they finish?

Simple Pipeline Approach



- Pipelining doesn't help latency of single task, it helps throughput of entire workload
- Pipeline rate limited by slowest pipeline stage
- Unbalanced lengths of pipeline stages reduce speedup
- Potential speedup = Number of pipeline stages
- Time to feed and bleed pipeline reduces speedup

Pipelined laundry takes 3.5 hours for 4 loads



A form of deadlock:

when two or more processes continue to execute, but make no progress towards the ultimate goal.

Examples:

- When there is an endless loop in program execution: e.g. it occurs when a process repeats itself, because it continues to receive erroneous information.
- When a process that calls another process is itself called by that process, and there is no logic to detect this situation and stop the operation.

So a livelock differs from a "deadlock," in that processing continues to take place, rather than just waiting in an idle loop.

Livelock: Basic XC examples

LIVELOCK



```
while (x) {  
  ch1 :=> x  
  ... Do something  
  ch2 <: y  
}
```

```
while (x) (  
  ... Do something  
  ch1 <: x  
  ch2 :=> y  
)
```

Two processes that will infinitely just communicate amongst themselves while X never becomes 0.

Basic Disjointness Rules [for Variables]

The rules for disjointness on a set of threads $T_0 \dots T_i$ and a set of variables $V_0 \dots V_j$:

- ❖ If thread T_x contains any modification to variable V_m then none of the other threads ($T_t; t \neq x$) are allowed to use V_m
- ❖ If thread T_x contains a reference to variable V_m then none of the other threads ($T_t; t \neq x$) are allowed to modify V_m
- ❖ If thread T_x contains a reference to port V_p then none of the other threads are allowed to use V_p
- a group of threads can have **shared read-only access** to a local variable, but only a single thread can have **exclusive read-write access**.
- disjointness rules for variables guarantee that each thread has a well-defined meaning that is independent of the order in which instructions in other threads are scheduled

Examples

```
main (void) {  
    int i=1, j=2, k=3;  
    par {  
        i = k + 1;           // Thread X  
        j = k - 1;           // Thread Y  
    }  
}
```

Legal or not?

```
main (void) {  
    int i=1, j=2, k;  
    par {  
        i = j + 1;           // Thread X  
        k = i - 1;           // Thread Y  
    }  
}
```

Legal or not?

```
main (void) {  
    int a[2];  
    par {  
        a[0] = f(0);         // Thread X  
        a[1] = f(1);         // Thread Y  
    }  
}
```

Legal or not?