

COMS20001 - Concurrent Computing

www.cs.bris.ac.uk/Teaching/Resources/COMS20001

Lecture 14

CSP: Liveness, Deadlock, Livelock



Sion Hannuna | hannuna@cs.bris.ac.uk
Tilo Burghardt | tilo@cs.bris.ac.uk
Dan Page | daniel.page@bristol.ac.uk


Recap: Refusals and Failures

We write P/tr for the process whose behaviour is whatever P could do after the trace tr has been observed.

Failures of a process:

$$failures(P) = \{(tr, X) \mid tr \in traces(P) \text{ and } X \in refusals(P/tr)\}$$

■ $P = a \rightarrow b \rightarrow STOP$ with $\alpha(P) = \{a, b\}$



$$traces(P) = \{\langle \rangle, \langle a \rangle, \langle a, b \rangle\}$$

$$refusals(P/\langle \rangle) = \{\{\}, \{b\}\}$$

$$refusals(P/\langle a \rangle) = \{\{\}, \{a\}\}$$

$$refusals(P/\langle a, b \rangle) = \{\{\}, \{a\}, \{b\}, \{a, b\}\}$$

$$failures(P) = \{(\langle \rangle, \{\}), (\langle \rangle, \{b\}), (\langle a \rangle, \{\}), (\langle a \rangle, \{a\}), (\langle a, b \rangle, \{\}), (\langle a, b \rangle, \{a\}), (\langle a, b \rangle, \{b\}), (\langle a, b \rangle, \{a, b\})\}$$

Failure Refinement I

Failure refinement is defined in a similar way to trace refinement:

$$\boxed{P \sqsubseteq_F Q \text{ if and only if } failures(Q) \subseteq failures(P)}$$

(Pronounce: " P is failure refined by Q ")

Failure refinement in specifications:

- $SPEC = a \rightarrow b \rightarrow SPEC$

-  Use $SPEC$ with trace refinement, get a *safety* specification!

- Find some processes P which satisfy $SPEC \sqsubseteq_T P$.

$P = STOP, P = a \rightarrow STOP, P = a \rightarrow b \rightarrow STOP, \dots$

- What effect has $SPEC \sqsubseteq_F P$?  First, calculate $failures(SPEC)$!

Failure Refinement II

$$\begin{aligned} \text{failures}(\text{SPEC}) = & \{(\langle a, b \rangle^n \frown \langle a \rangle, \emptyset) \mid n \geq 0\} \\ & \cup \{(\langle a, b \rangle^n \frown \langle a \rangle, \{a\}) \mid n \geq 0\} \\ & \cup \{(\langle a, b \rangle^n, \emptyset) \mid n \geq 0\} \\ & \cup \{(\langle a, b \rangle^n, \{b\}) \mid n \geq 0\} \end{aligned}$$

To find out whether $\text{SPEC} \sqsubseteq_F \text{STOP}$, calculate:

$$\text{failures}(\text{STOP}) = \{(\langle \rangle, \emptyset), (\langle \rangle, \{a\}), (\langle \rangle, \{b\}), (\langle \rangle, \{a, b\})\}$$

**Pairs $(\langle \rangle, \{a\})$ and $(\langle \rangle, \{a, b\})$ are failures of STOP , but not of SPEC .
Hence, $\text{SPEC} \not\sqsubseteq_F \text{STOP}$.**

Now, consider $P = a \rightarrow \text{STOP}$.

$$\text{failures}(P) = \{(\langle \rangle, \emptyset), (\langle \rangle, \{b\}), (\langle a \rangle, \emptyset), (\langle a \rangle, \{a\}), (\langle a \rangle, \{b\}), (\langle a \rangle, \{a, b\})\}$$

Failure pairs $(\langle a \rangle, \{b\})$ and $(\langle a \rangle, \{a, b\})$ are failures of P but not of SPEC ; so again $\text{SPEC} \not\sqsubseteq_F P$.

Liveness (guaranteed execution of some behaviour)

$SPEC \sqsubseteq_F P$ is a *liveness* specification which requires P to do certain events.

■ Which definitions of P satisfy $SPEC = a \rightarrow b \rightarrow SPEC$?

Obviously $P = a \rightarrow b \rightarrow P$ does.

💡 It is (in this case) the only process satisfying this specification!

(Specification is too tight; pins down implementation precisely.)

Liveness Specification Example & Hiding

Process P with alphabet $\{a, b, c\}$.

- Want to specify that P must be able to do an infinite sequence of alternating a and b events, starting with a .
- We do not care about c events.

■ Use process $ALT = a \rightarrow b \rightarrow ALT$ as before.

Allow c events to occur freely through hiding: $ALT \sqsubseteq_F (P \setminus \{c\})$

■ Definitions of P satisfying this specification include: $P = a \rightarrow b \rightarrow P$,
 $P = c \rightarrow a \rightarrow c \rightarrow c \rightarrow b \rightarrow P$, $P = a \rightarrow b \rightarrow c \rightarrow P$.

💡 All are the same as ALT when c is hidden!

■ Definitions of P not satisfying this specification include: $P = STOP$,
 $P = a \rightarrow b \rightarrow (P \sqcap a \rightarrow c \rightarrow STOP)$

Safety Spec vs. Liveness Spec

Saying that $tr \in traces(P)$ is a *positive* statement.

💡 Describes something that P can do!

$SPEC \sqsubseteq_T P$ puts limit on traces that P can do; restricts behaviour.

💡 P may fail a *safety* (trace) specification by doing too much.

Saying that $(tr, X) \in failures(P)$ is a *negative* statement.

💡 Describes something that P cannot do!

$SPEC \sqsubseteq_F P$ puts limit on what P can fail to do.

⇒ Requires P to accept at least a certain range of behaviours.

💡 P may fail a *liveness* (failure) specification by refusing too much, i.e. by not doing enough.

Example: Moving Furniture

- Two furniture movers need to move a table and a piano. Each requires two people to lift it.

$$\begin{aligned} PETE &= lift.piano \rightarrow PETE \\ &\sqcap lift.table \rightarrow PETE \end{aligned}$$

$$\begin{aligned} DAVE &= lift.piano \rightarrow DAVE \\ &\sqcap lift.table \rightarrow DAVE \end{aligned}$$

$$TEAM = PETE \parallel DAVE$$

- 💡 Both Pete and Dave make their decisions independently! (□)
 - If both make same choice, they can cooperate in moving an object.

Deadlock Example: Moving Furniture

- If their choices are different, ...

$$PETE \xrightarrow{\tau} lift.piano \rightarrow PETE$$
$$DAVE \xrightarrow{\tau} lift.table \rightarrow DAVE$$

💡 $lift.piano \rightarrow PETE \parallel lift.table \rightarrow DAVE$ cannot do anything.
(It is equivalent to the process *STOP*!)

A state of a process is **deadlocked** if it can **refuse** to do every event.
STOP is the simplest deadlocked process.

Deadlock Example: Children Painting

■ Ella and Kate share a paint box and an easel.

$$ELLA = ella.get.box \rightarrow ella.get.easel \rightarrow ella.paint \rightarrow \\ ella.put.box \rightarrow ella.put.easel \rightarrow ELLA$$
$$KATE = kate.get.easel \rightarrow kate.get.box \rightarrow kate.paint \rightarrow \\ kate.put.easel \rightarrow kate.put.box \rightarrow KATE$$
$$EASEL = ella.get.easel \rightarrow ella.put.easel \rightarrow EASEL \\ \square kate.get.easel \rightarrow kate.put.easel \rightarrow EASEL$$
$$BOX = ella.get.box \rightarrow ella.put.box \rightarrow BOX \\ \square kate.get.box \rightarrow kate.put.box \rightarrow BOX$$

Combination of two children, box and easel:

$$PAINTING = ELLA \parallel KATE \parallel EASEL \parallel BOX$$

(Assume synchronisation on (intersection of) individual alphabets.)

Conditions for Deadlock

Coffman, Elphick and Shoshani identified 4 *necessary and sufficient* conditions for deadlock [System Deadlocks. ACM Computing Surveys 3, 2 (June), p. 67-78, 1971.]

1. Agents claim exclusive control of the resources they require.

⇒ **“Mutual exclusion”** condition

2. Agents hold resources already allocated to them while waiting for additional resources.

⇒ **“Wait for”** condition

3. Resources cannot be forcibly removed from the agent holding them until the resources are used to completion.

⇒ **“No preemption”** condition

4. A circular chain of agents exists, s.t. each agent holds one or more resources that are being requested by the next task in the chain.

⇒ **“Circular wait”** condition

Breaking Deadlock

Aim: System in which possibility of deadlock is excluded a priori.

💡 Ensure that at least one of the conditions is not satisfied!

⇒ Constrain the way in which requests for resources are made.

- Usually “**Mutual exclusion**” condition **cannot be denied**.
- Each agent must request all its required resources at once and cannot proceed until all have been granted. 💡 Make it **atomic**!
⇒ “**Wait for**” condition **denied**.
- If an agent holding certain resources is denied a further request, that agent must release its original resources and, if necessary, request them again together with the original resources.
⇒ “**No preemption**” condition **denied**.
- Imposition of a *linear* ordering of resource types.
⇒ “**Circular wait**” condition **denied**.

Example: Breaking Deadlock with Semaphore

💡 Introduce a semaphore to break the deadlock!

■ Ella and Kate share a paint box and an easel. (SIMPLIFIED)

$ELLA = ella.getsem \rightarrow ella.get.box \rightarrow ella.get.easel \rightarrow ella.paint \rightarrow$
 $ella.put.box \rightarrow ella.put.easel \rightarrow ella.putsem \rightarrow ELLA$

$KATE = kate.getsem \rightarrow kate.get.easel \rightarrow kate.get.box \rightarrow kate.paint \rightarrow$
 $kate.put.easel \rightarrow kate.put.box \rightarrow kate.putsem \rightarrow KATE$

$BSEM' = ella.getsem \rightarrow ella.putsem \rightarrow BSEM'$

□ $kate.getsem \rightarrow kate.putsem \rightarrow BSEM'$

💡 To achieve desired synchronisation with semaphore (in CSP), we need one channel for each process that uses the semaphore.

Now deadlock-free: $PAINTING = ELLA || KATE || BSEM || EASEL || BOX$

(Assume $||$ abbreviates alphabetised parallel here, so synchronisation is on intersection of individual alphabets.)

Example: Breaking Deadlock via Preemption

💡 If an agent holding certain resources is denied a further request, that agent must release its original resources and, if necessary, request them again together with the original resources.

⇒ “No preemption” condition **denied**.

■ Let Ella return tools before they have been used, rather than wait indefinitely for them to be available.

$$\begin{aligned} ELLA' &= ella.get.box \rightarrow (ella.put.box \rightarrow ELLA' \\ &\quad \square ella.get.easel \rightarrow ella.paint \rightarrow \\ &\quad ella.put.box \rightarrow ella.put.easel \rightarrow ELLA') \\ &\quad \square ella.get.easel \rightarrow (ella.put.easel \rightarrow ELLA' \\ &\quad \quad \square ella.get.box \rightarrow ella.paint \rightarrow \\ &\quad \quad ella.put.easel \rightarrow ella.put.box \rightarrow ELLA') \end{aligned}$$

Will this get rid of the deadlock? **Yes, but it introduces a “livelock”!**

Livelock ... liveness without progress

Difficult to distinguish *wanted* and *unwanted* repeated events!

💡 Hide "auxiliary" events, just focus on important events!

■ In Ella/Kate example we are only interested in *paint* events:

$SYSTEM = PAINTING \setminus \{ella, kate\} . \{put, get\} . \{easel, box\}$

where $PAINTING = ELLA' || KATE || EASEL || BOX$

In CSP, possibility of an infinite sequence of τ events is called livelock or divergence.

💡 A divergent process cannot be guaranteed to make any progress.

■ Ella can loop forever without achieving any painting!

■ Busy waiting (e.g. in mutual exclusion algorithms) counts as divergence!

Failure Refinement cannot detect Livelock!

Trace model: $Spec \sqsubseteq_T Imp$ iff $traces(Spec) \supseteq traces(Imp)$

Failure model: $Spec \sqsubseteq_F Imp$ iff $failures(Spec) \supseteq failures(Imp)$


- Failures model does not allow to detect whether a process might livelock!

■ $SPEC \sqsubseteq_T IMP \setminus \{b\}$ and $SPEC \sqsubseteq_F IMP \setminus \{b\}$ where

$$SPEC = a \rightarrow SPEC$$

$$IMP = a \rightarrow IMP \sqcap b \rightarrow IMP$$

Both refinements hold. BUT IMP livelocks!



Remember:

- The process $P \setminus A$ undergoes same executions as P , but events from A occur as *internal events* τ , which are *not visible*.



Failure-Divergence Refinement

We need a model that is more thorough than stable failures model!

💡 **Failures Divergences model** represents process by its stable failures and its divergences.

A *divergence* is a finite trace during or after which the process can perform an infinite sequence of consecutive internal events.

$Spec \sqsubseteq_{FD} Imp$ iff
 $failures(Spec) \supseteq failures(Imp)$ and $divergences(Spec) \supseteq divergences(Imp)$



Summary of CSP...

Traces model:

- Safety properties (do no wrong).
- The sequences of traces that a process can perform.

Failures model:

- Liveness properties (do something right).
- Deadlock freedom.
- Pairs of traces and the refusals that may occur after them.

Failures-Divergence model:

- Livelock freedom.
- Failures plus the traces that lead to divergence.