

► Problem:

- memory represents short-term (transient, volatile) storage, so
- we need explicit control over long-term (permanent, non-volatile) storage, st.
 1. **persistence** : decouple lifetime of data from any given process
 2. **organisation** : unambiguously and conveniently identify data
 3. **efficiency** : low-latency, random access to data
 4. **robustness** : accessing or duplicating data is error-free

► Solution: we need to consider

- what an appropriate system call interface should be, *and*
- how to map the semantics of said interface onto one or more concrete storage devices.

Concept (1)

Definition

- ▶ A **file** is a logical unit of (stored) data.
- ▶ A **file system** is an abstraction mechanism: it allows logical manipulation of files, without knowledge of their physical representation.

Concept (1)

Definition

A **file system** provides a mapping

$\text{identifier} \mapsto (\text{meta-data}, \text{data}),$

plus a mechanism to manage (concurrent) manipulation of both

data	\simeq	content
meta-data	\simeq	structure

Concept (1)

Definition

A **file system** provides a mapping

identifier \mapsto (meta-data, data),

plus a mechanism to manage (concurrent) manipulation of both

data	\simeq	content
meta-data	\simeq	structure

- **Question:** why be so abstract?

Definition

A **file system** provides a mapping

identifier \mapsto (meta-data, data),


plus a mechanism to manage (concurrent) manipulation of both

data	\simeq	content
meta-data	\simeq	structure

- ▶ **Question:** why be so abstract?
- ▶ **Answer:** file systems support multiple use-cases, e.g.,
 1. general-purpose data storage,
 2. special-purpose data storage (e.g., swap space [4]), or
 3. interface with kernelso saying “file” rather than “data” *may* be artificially limiting.

Concept (2)

- **Option**: the mapping can range between



root directory

```
|— foo.txt  
|— bar.txt  
|— baz.txt  
|— ...
```


with $l > 1$ implying

- entries may be **directories**,
- the identifier specifying an entry includes a (potentially implicit) **path**,
- paths can be
 - **absolute** (from *root* directory) or
 - **relative** (from *some* directory)

and hence needn't be unique.

Concept (2)

- **Option**: the mapping can range between




with $l > 1$ implying

- entries may be **directories**,
- the identifier specifying an entry includes a (potentially implicit) **path**,
- paths can be
 - **absolute** (from *root* directory) or
 - **relative** (from *some* directory)

and hence needn't be unique.

root directory




Concept (2)

► **Option:** the root can be

1. a **volume** identifier, or
2. a directory

where

- the former implies multiple, segregated hierarchies,
- the latter implies one, unified hierarchy ...
- ... we **mount** a volume at a **mount point**.



Concept (2)


► **Option:** entries in the hierarchy include

- a file,
- a directory,
- a **symbolic link**, and
- a **device node**

where

- entry types may be differentiated via meta-data or embedded magic numbers,
- links are often categorised as either
 - **hard**, or
 - **soft**

but, either way, imply the hierarchy is now a (acyclic) *graph* (vs. a tree).



Mechanism: POSIX(ish) system call interface (1)

- **Assumption:** underlying file system allows us to view data as a byte sequence, i.e.,

$$\text{identifier} \mapsto (\text{meta-data}, \text{data}) = (\text{meta-data}, \boxed{\begin{array}{|c|c|c|c|} \hline d_0 & d_1 & \cdots & d_{n-1} \\ \hline \end{array}})$$

↑
read/write pointer

that supports

1. automatic extensibility, and
2. random access (via **seek** operations).

Mechanism: POSIX(ish) system call interface (2)

- ▶ Based on this assumption, the kernel must (at least)
 1. maintain a global **mount table** that captures the hierarchy,

Mechanism: POSIX(ish) system call interface (2)

- ▶ Based on this assumption, the kernel must (at least)
 2. maintain a per process **file descriptor table** that captures
 - ▶ a **File Control Block (FCB)** of physical addressing information,
 - ▶ the mode the entry was opened in (e.g., read, write, or read/write),
 - ▶ the current read/write pointerthat indexes into a global **file table** tracking open entries,

Mechanism: POSIX(ish) system call interface (2)

- Based on this assumption, the kernel must (at least)

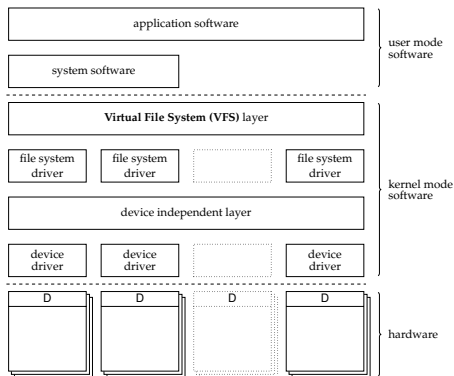
3. support a suite of system calls, e.g.,

Function	Reference	Purpose
creat	[11, Page 702]	create a file
open	[11, Page 1379]	open a file
close	[11, Page 676]	close a file
unlink	[11, Page 2154]	delete a file
write	[11, Page 2263]	write to a file
read	[11, Page 1737]	read from a file
lseek	[11, Page 1265]	move read/write pointer

which are ...

Mechanism: POSIX(ish) system call interface (3)

- ▶ ... (typically) exposed via a **Virtual File System (VFS)** layer

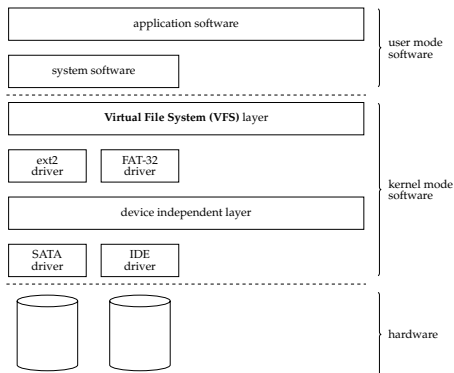


offering

1. a uniform interface to
 - ▶ multiple heterogeneous concrete file systems, and
 - ▶ “device-less” pseudo-files
- plus
2. various optimisation and translation operations.

Mechanism: POSIX(ish) system call interface (3)

- ▶ ... (typically) exposed via a **Virtual File System (VFS)** layer

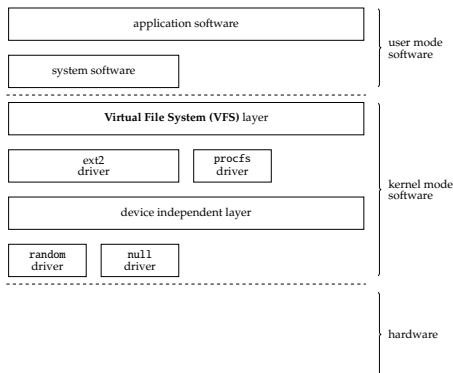


offering

1. a uniform interface to
 - ▶ multiple heterogeneous concrete file systems, and
 - ▶ “device-less” pseudo-files
- plus
2. various optimisation and translation operations.

Mechanism: POSIX(ish) system call interface (3)

- ▶ ... (typically) exposed via a **Virtual File System (VFS)** layer

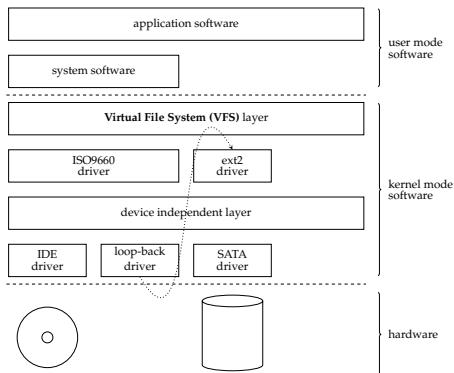


offering

1. a uniform interface to
 - ▶ multiple heterogeneous concrete file systems, and
 - ▶ “device-less” pseudo-files
- plus
2. various optimisation and translation operations.

Mechanism: POSIX(ish) system call interface (3)

- ... (typically) exposed via a **Virtual File System (VFS)** layer

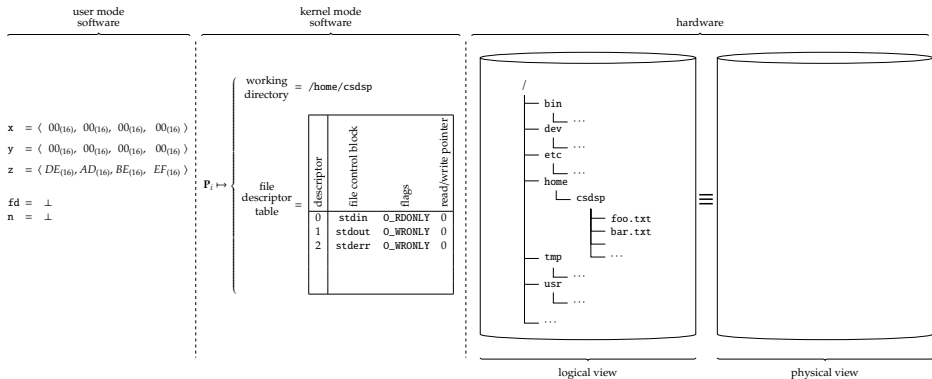


offering

1. a uniform interface to
 - multiple heterogeneous concrete file systems, and
 - “device-less” pseudo-files
2. various optimisation and translation operations.

Mechanism: POSIX(ish) system call interface (4)

► Example:

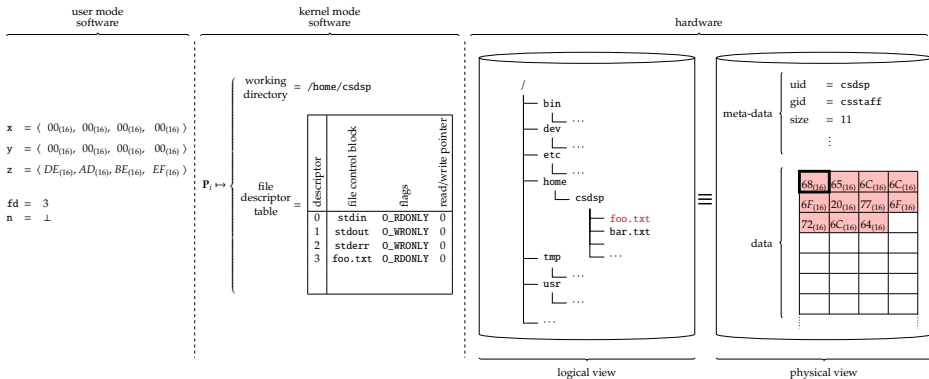


Mechanism: POSIX(ish) system call interface (4)

- **Example:** if the user mode process executes

```
fd = open( "foo.txt", O_RDONLY )
```

then the result is described by

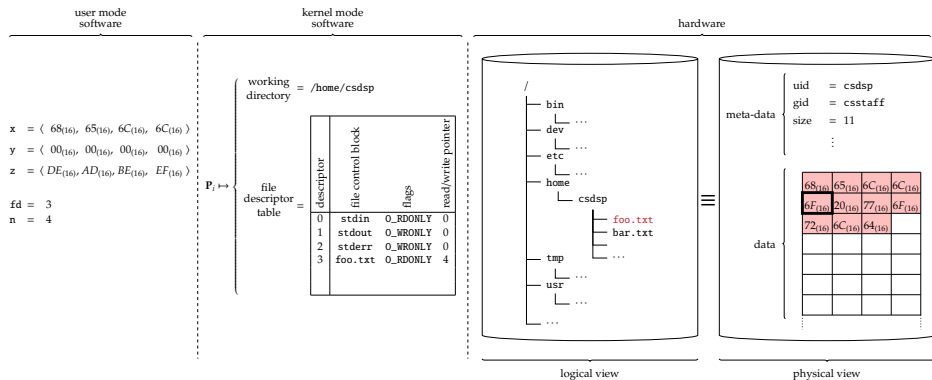


Mechanism: POSIX(ish) system call interface (4)

- **Example:** if the user mode process executes

```
n = read( fd, x, 4 )
```

then the result is described by

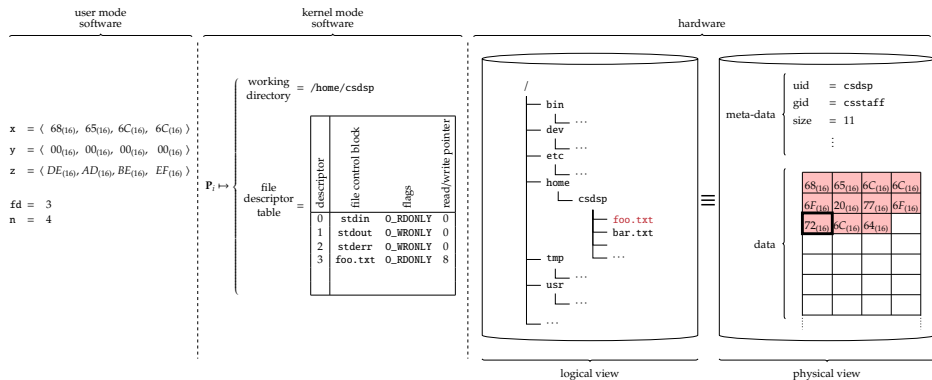


Mechanism: POSIX(ish) system call interface (4)

- **Example:** if the user mode process executes

`lseek(fd, +8, SEEK_SET)`

then the result is described by

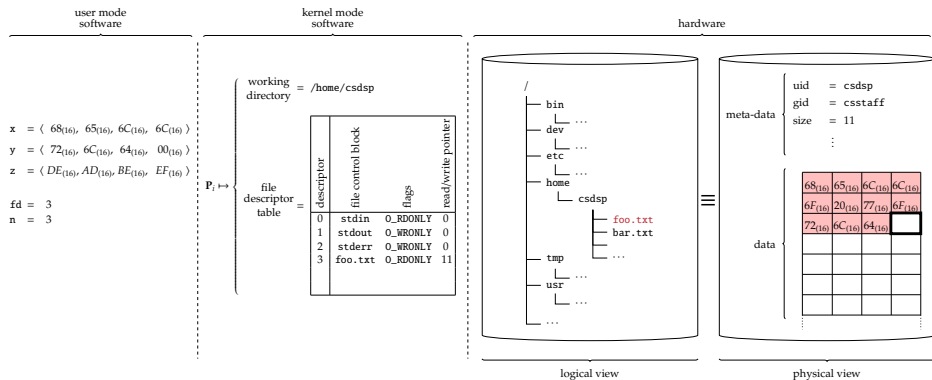


Mechanism: POSIX(ish) system call interface (4)

- **Example:** if the user mode process executes

```
n = read( fd, y, 4 )
```

then the result is described by

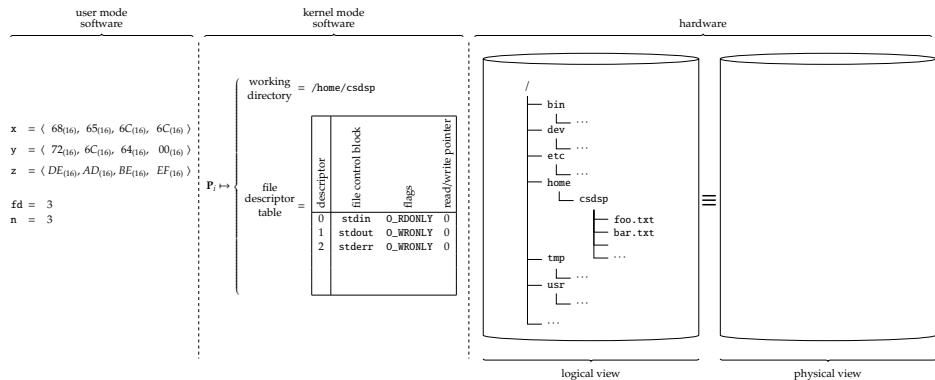


Mechanism: POSIX(ish) system call interface (4)

- **Example:** if the user mode process executes

`close(fd)`

then the result is described by

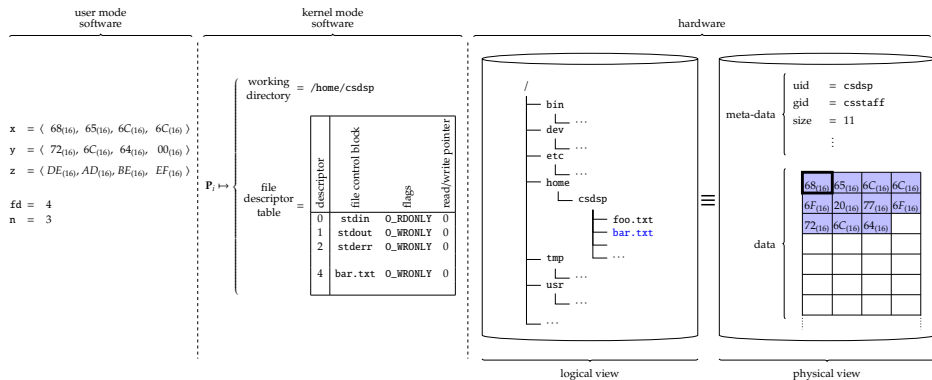


Mechanism: POSIX(ish) system call interface (4)

- **Example:** if the user mode process executes

```
fd = open( "bar.txt", O_WRONLY )
```

then the result is described by

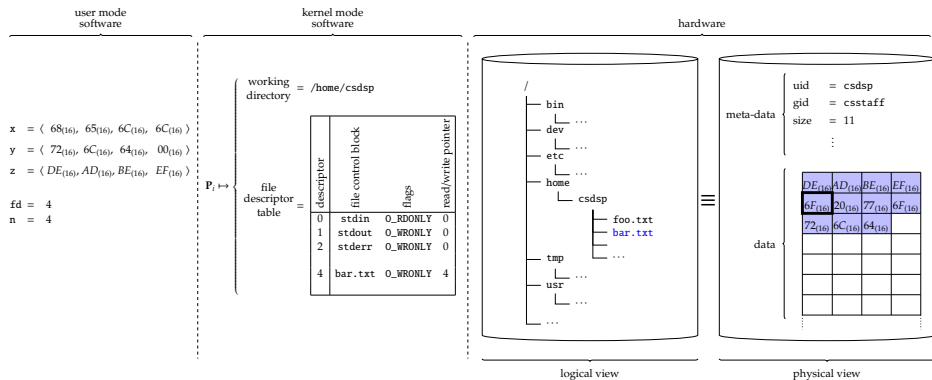


Mechanism: POSIX(ish) system call interface (4)

- **Example:** if the user mode process executes

```
n = write( fd, z, 4 )
```

then the result is described by

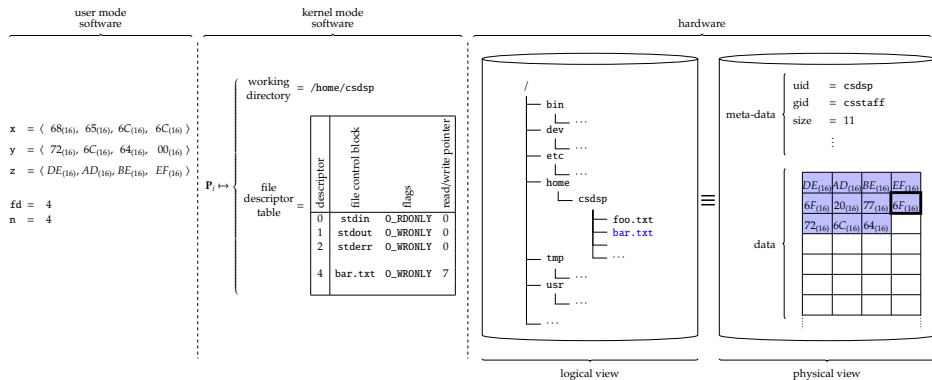


Mechanism: POSIX(ish) system call interface (4)

- **Example:** if the user mode process executes

`lseek(fd, -4, SEEK_END)`

then the result is described by

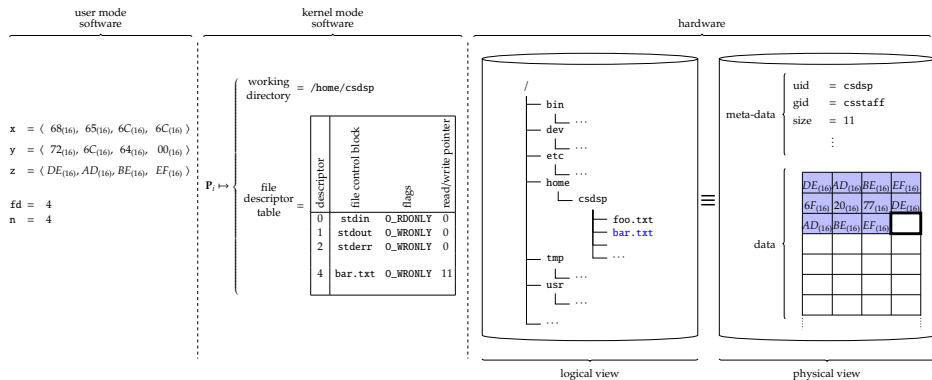


Mechanism: POSIX(ish) system call interface (4)

- **Example:** if the user mode process executes

```
n = write( fd, z, 4 )
```

then the result is described by

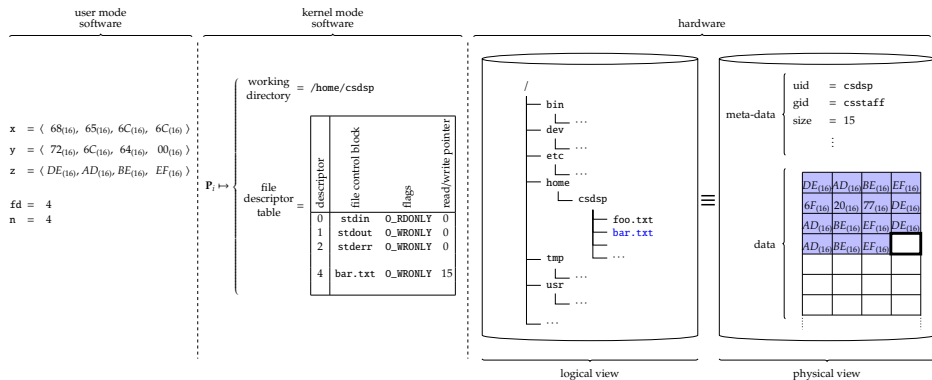


Mechanism: POSIX(ish) system call interface (4)

- **Example:** if the user mode process executes

```
n = write( fd, z, 4 )
```

then the result is described by

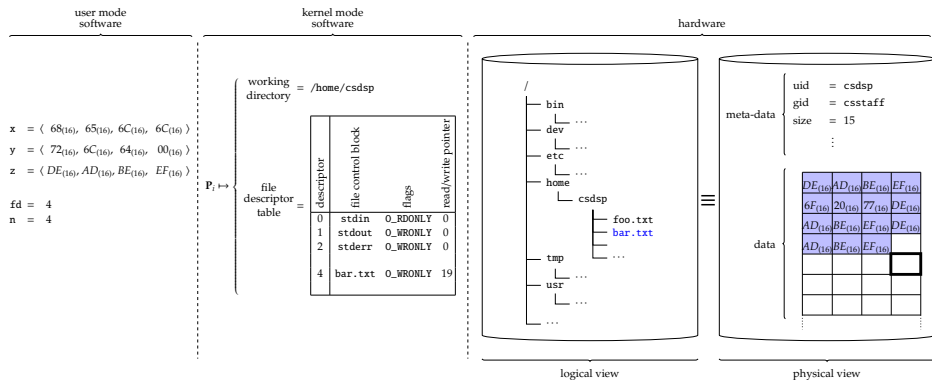


Mechanism: POSIX(ish) system call interface (4)

- **Example:** if the user mode process executes

`lseek(fd, +4, SEEK_END)`

then the result is described by

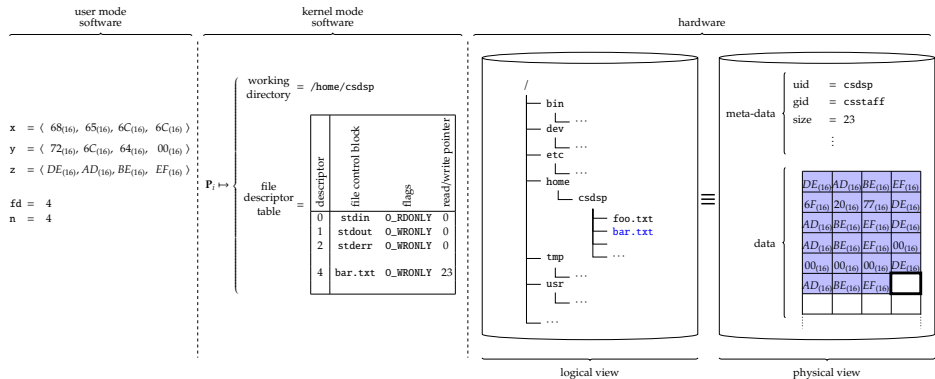


Mechanism: POSIX(ish) system call interface (4)

- **Example:** if the user mode process executes

```
n = write( fd, z, 4 )
```

then the result is described by

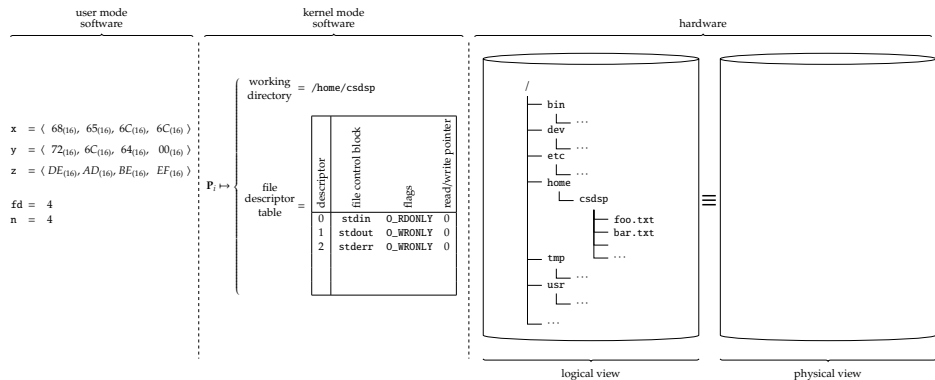


Mechanism: POSIX(ish) system call interface (4)

- **Example:** if the user mode process executes

`close(fd)`

then the result is described by

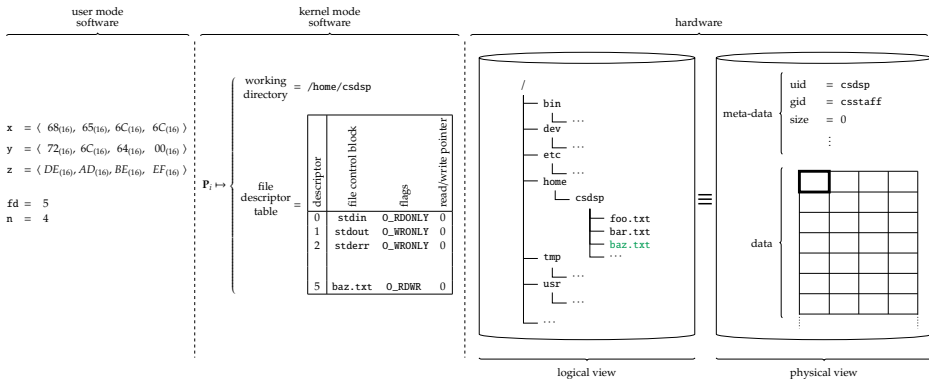


Mechanism: POSIX(ish) system call interface (4)

- **Example:** if the user mode process executes

```
fd = creat( "baz.txt", O_WRONLY )
```

then the result is described by

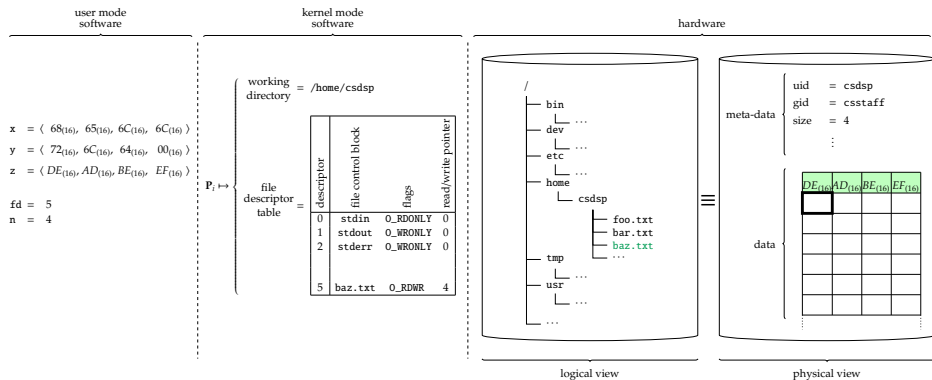


Mechanism: POSIX(ish) system call interface (4)

- **Example:** if the user mode process executes

```
n = write( fd, z, 4 )
```

then the result is described by

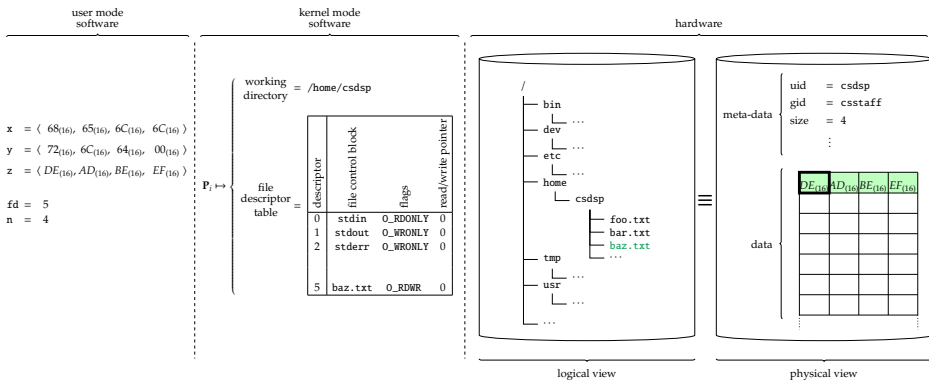


Mechanism: POSIX(ish) system call interface (4)

- **Example:** if the user mode process executes

`lseek(fd, +0, SEEK_SET)`

then the result is described by

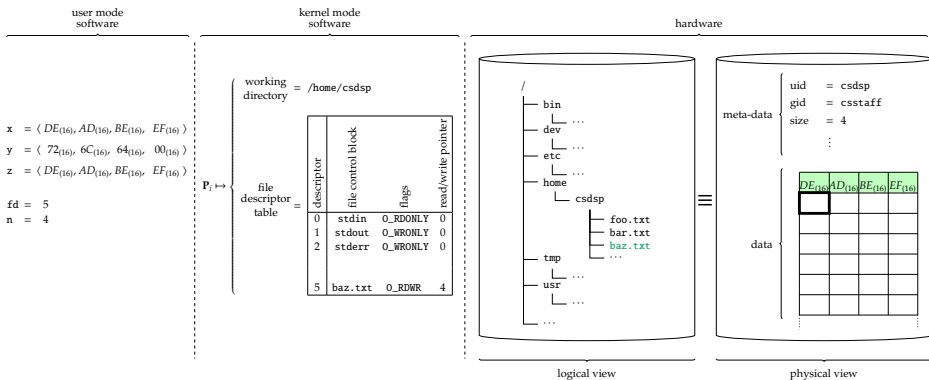


Mechanism: POSIX(ish) system call interface (4)

- **Example:** if the user mode process executes

```
n = read( fd, x, 4 )
```

then the result is described by

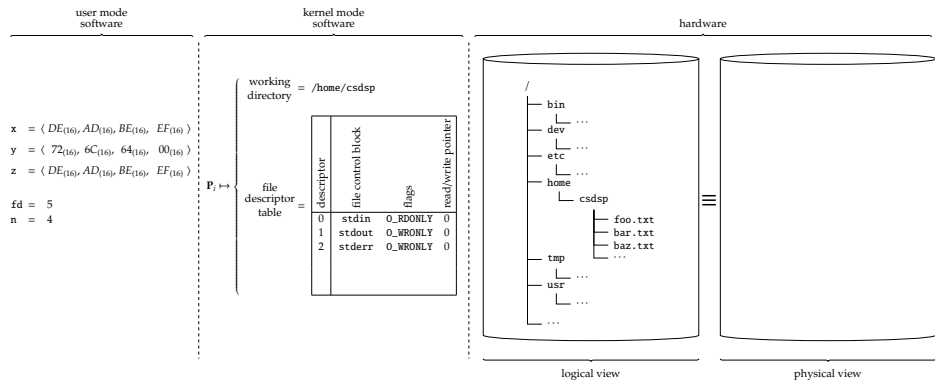


Mechanism: POSIX(ish) system call interface (4)

- ▶ **Example:** if the user mode process executes

`close(fd)`

then the result is described by

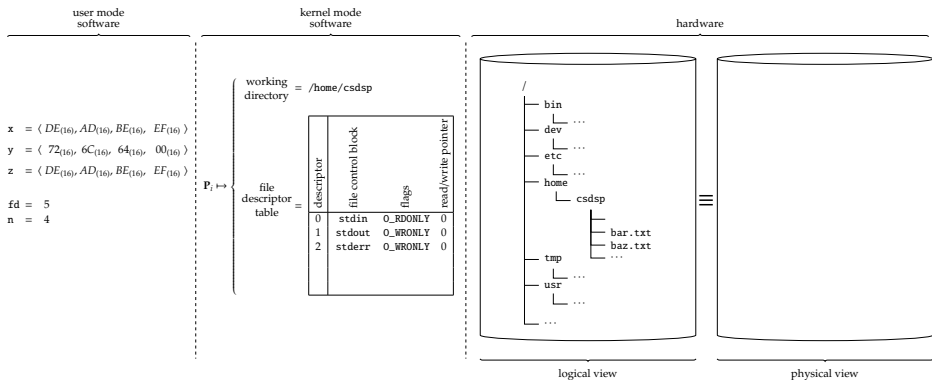


Mechanism: POSIX(ish) system call interface (4)

- **Example:** if the user mode process executes

```
unlink( "foo.txt" )
```

then the result is described by



Concept (1)

- ▶ ... so far so good, *but* we need to explore
 1. how the file system supports our assumed access model, *and*
 2. how the underling **storage device** supports the file system.


Concept (2): (mass storage) devices \leadsto blocks

- ▶ We add structure to the medium via several steps

medium =

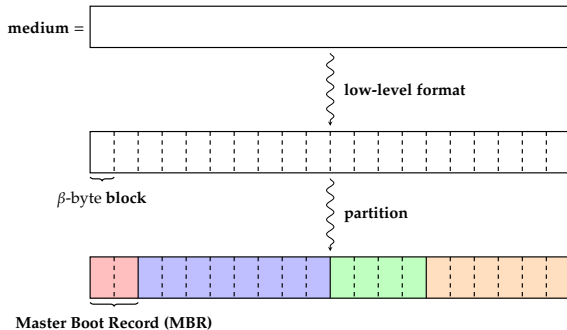
Concept (2): (mass storage) devices \leadsto blocks

- We add structure to the medium via several steps



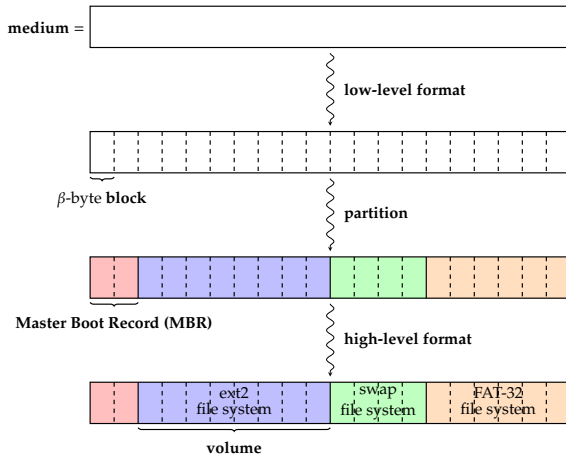
Concept (2): (mass storage) devices \leadsto blocks

- We add structure to the medium via several steps



Concept (2): (mass storage) devices \leadsto blocks


- ▶ We add structure to the medium via several steps



then ...


Concept (3): (mass storage) devices \leadsto blocks

- ... assume an interface as previously described, e.g.



Concept (3): (mass storage) devices \leadsto blocks

- ▶ ... assume an interface as previously described, e.g.



but, since efficiency is crucial we will (typically)

- ▶ amortise overhead by fixing transferring β -byte blocks, and
- ▶ buffer and/or cache accesses.

Concept (4): blocks \leadsto files

► **Challenge:** given a device with

- fixed number of logical blocks and
- fixed sized logical blocks,

realise (hierarchical) file system supporting


- representation and
- manipulation

of

- variable number of files, and
- variable sized files.


► **Solution:** we need

1. an allocation algorithm, and
2. a data structure to capture the current allocation state.



Concept (4): blocks \leadsto files

- ▶ **Challenge:** given a device with
 - ▶ fixed number of logical blocks and
 - ▶ fixed sized logical blocks,realise (hierarchical) file system supporting
 - ▶ representation and
 - ▶ manipulationof
 - ▶ variable number of files, and
 - ▶ variable sized files.
- ▶ **Idea: contiguous allocation.**
 - allocation is more challenging,
 - + sequential *and* random access is efficient,
 - internal *and* external fragmentation,
 - + no storage overhead.



Concept (4): blocks \leadsto files

► Challenge: given a device with

- fixed number of logical blocks and
- fixed sized logical blocks,

realise (hierarchical) file system supporting


- representation and
- manipulation

of

- variable number of files, and
- variable sized files.

► Idea: **linked allocation.**

- + allocation is less challenging,
- + sequential access is efficient,
- random access is inefficient,
- + internal fragmentation only,
- some storage overhead due to pointers.



Concept (4): blocks \leadsto files

► Challenge: given a device with

- fixed number of logical blocks and
- fixed sized logical blocks,

realise (hierarchical) file system supporting


- representation and
- manipulation

of

- variable number of files, and
- variable sized files.


► Idea: **linked allocation.**

- + allocation is less challenging,
- + sequential access is efficient,
- random access is inefficient,
- + internal fragmentation only,
- some storage overhead due to pointers.




Concept (4): blocks \leadsto files

- ▶ **Challenge:** given a device with
 - ▶ fixed number of logical blocks and
 - ▶ fixed sized logical blocks,realise (hierarchical) file system supporting
 - ▶ representation and
 - ▶ manipulationof
 - ▶ variable number of files, and
 - ▶ variable sized files.
- ▶ **Idea: indexed allocation.**
 - + allocation is less challenging,
 - + sequential *and* random access is efficient,
 - + internal fragmentation only,
 - some storage overhead due to pointers.



Concept (4): blocks \leadsto files

- ▶ **Challenge:** given a device with
 - ▶ fixed number of logical blocks and
 - ▶ fixed sized logical blocks,realise (hierarchical) file system supporting
 - ▶ representation and
 - ▶ manipulationof
 - ▶ variable number of files, and
 - ▶ variable sized files.
- ▶ **Problem:** what if the index block is full?
- ▶ **Solution(s):** use *multiple* index blocks, e.g., via
 1. linked list,
 2. linked tree, or
 3. various hybrid(s) ...



Concept (5): blocks \leadsto files

► **Problem:** larger storage capacity means more logical blocks, so

1. larger logical block addresses,
2. decreased access locality, and
3. greater overhead (in time and space) wrt. allocation.

► **Solution:** use a hybrid part contiguous, part non-contiguous approach, e.g.,

1. **cluster** \simeq fixed size, contiguous group of logical blocks:

► e.g., divide w -bit logical block address into two

$$\text{logical block group address} = \begin{array}{|c|c|} \hline \overset{w-1}{\text{offset}} & \overset{t}{\text{00...0}} \overset{t-1}{\text{0}} \\ \hline \end{array}$$

► each offset now addresses a contiguous group of 2^t logical blocks.

2. **extent** \simeq variable size, contiguous group of logical blocks:

► e.g., divide w -bit logical block address into two

$$\text{logical block group address} = \begin{array}{|c|c|} \hline \overset{w-1}{\text{offset}} & \overset{t}{\text{length}} \overset{t-1}{\text{0}} \\ \hline \end{array}$$

- each offset now addresses a contiguous group of logical blocks whose length is given by the t LSBs,
- this is more flexible, *but* yields complications wrt. seeking and allocation.

although from here on we ignore this option.

Concept (6): blocks \leadsto files

- **Problem:** each write requires one or more of

1. update the allocation state,
2. update the file meta-data, *and*
3. update the file data

which *must* be **atomic**: if not, the file system can become inconsistent.

- **Solution:**

- describe update in write-ahead log (or journal),
- commit update to file system iff. write to log is complete.

Concept (7): blocks \leadsto files

- ▶ **Question:** what *is* a directory?
- ▶ **Answer:** a mapping, e.g.,

identifier \mapsto (meta-data, data)

or

(identifier, meta-data) \mapsto data

which also hint at

1. options for where meta-data should reside, and
2. the fact a file might not *itself* have an identifier!

Concept (7): blocks \leadsto files

► **Question:** what *is* a directory?

► **Answer:** a mapping, e.g.,

identifier \mapsto (meta-data, data)

or

(identifier, meta-data) \mapsto data

which also hint at

1. options for where meta-data should reside, and
2. the fact a file might not *itself* have an identifier!

► **Problem:** *how* should a directory be represented?

► **Solution:**

1. list,
2. tree,
3. hash table,
4. ...

Concept (7): blocks \leadsto files

- ▶ **Question:** what *is* a directory?
- ▶ **Answer:** a mapping, e.g.,

identifier \mapsto (meta-data, data)

or

(identifier, meta-data) \mapsto data

which also hint at

1. options for where meta-data should reside, and
2. the fact a file might not *itself* have an identifier!

- ▶ **Problem:** *where* should a directory representation be stored?
- ▶ **Solution:**


1. as a file, plus special-purpose rules for access,
2. as a special-purpose structure,
3. ...

i.e., unified or segregated wrt. the rest of the file system.

Implementation (1): devices \leadsto file systems

Windows-centric: FAT

- **Idea: File Allocation Table (FAT) \approx fancy linked allocation.**



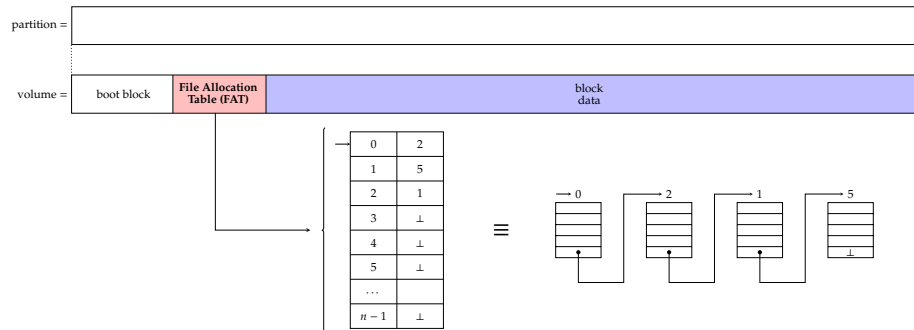
noting that

- + resolves issue of random access wrt. linked allocation,
- need to retain FAT in memory ... which, for n logical blocks, can be large!

Implementation (1): devices \leadsto file systems

Windows-centric: FAT

- **Idea: File Allocation Table (FAT) \approx fancy linked allocation.**




noting that

- + resolves issue of random access wrt. linked allocation,
- need to retain FAT in memory ... which, for n logical blocks, can be large!

Implementation (2): devices \leadsto file systems

UNIX-centric: UFS

- **Idea: Unix File System (UFS)** [14, Section 4] \approx fancy indexed allocation.



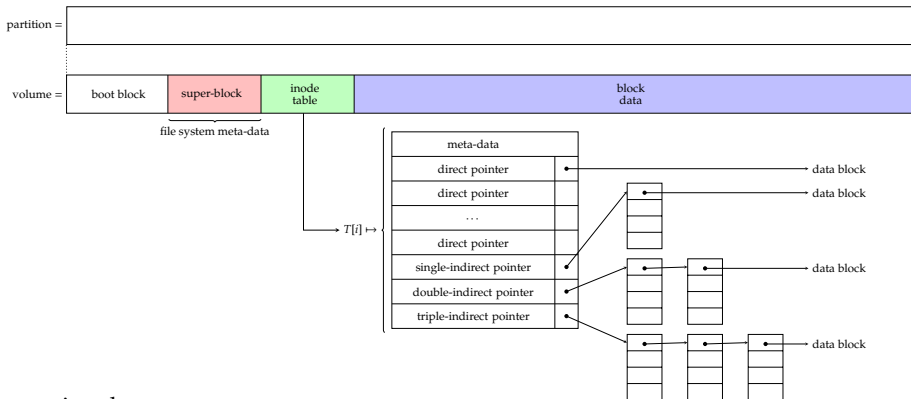
noting that

- + inodes are of (small) fixed size, so indexing into the inode table is efficient,
- linked representation of free space (for inodes *and* blocks).

Implementation (2): devices \leadsto file systems

UNIX-centric: UFS

- **Idea: Unix File System (UFS)** [14, Section 4] \approx fancy indexed allocation.




noting that

- + inodes are of (small) fixed size, so indexing into the inode table is efficient,
- linked representation of free space (for inodes *and* blocks).

Implementation (3): devices \leadsto file systems

UNIX-centric: FFS

- **Idea: Fast File System (FFS) [12]** \approx UFS + larger block size ($\geq 4\text{KiB}$).




noting that

- + bit-map representation of free space (for inodes *and* blocks),
- + redundant copy of super-block improves fault tolerance (for overhead of space),
- + cylinder groups increase access locality,
- + includes additional features such as soft links.

Implementation (4): devices \leadsto file systems

UNIX-centric: ext2

- **Idea: Second Extended File System (ext2) [3] \approx FFS + caching.**



noting that

- + improved directory representation via hash tables [3, Section 4.2],
- + caching and asynchronous writes improve performance ...
- ... but with disadvantages wrt. coherence (viz. robustness).

► Take away points:

- This is a broad and complex topic: it involves (at least)
 1. a hardware aspect:
 - an interrupt controller,
 - a block device
 2. a low(er)-level software aspect:
 - an interrupt handler,
 - a device driver,
 - a file system driver
 3. a high(er)-level software aspect:
 - some data structures (e.g., mount and file descriptor tables),
 - any relevant POSIX system calls (e.g., `write`)
- Keep in mind that, even then,
 - we've excluded and/or simplified various (sub-)topics,
 - there are numerous trade-offs involved, meaning it is often hard to identify one ideal solution.

Additional Reading

- ▶ *Wikipedia: File system*. URL: http://en.wikipedia.org/wiki/File_system.
- ▶ R. Love. “Chapter 2: File I/O”. In: *Linux System Programming*. 2nd ed. O’Reilly, 2013.
- ▶ R. Love. “Chapter 8: File and directory management”. In: *Linux System Programming*. 2nd ed. O’Reilly, 2013.
- ▶ A. Silberschatz, P.B. Galvin, and G. Gagne. “Chapter 10: File system”. In: *Operating System Concepts*. 9th ed. Wiley, 2014.
- ▶ A. Silberschatz, P.B. Galvin, and G. Gagne. “Chapter 11: Implementing file systems”. In: *Operating System Concepts*. 9th ed. Wiley, 2014.
- ▶ A. Silberschatz, P.B. Galvin, and G. Gagne. “Chapter 12: Mass storage structure”. In: *Operating System Concepts*. 9th ed. Wiley, 2014.
- ▶ A.S. Tanenbaum and H. Bos. “Chapter 4: File systems”. In: *Modern Operating Systems*. 4th ed. Pearson, 2015.

References

- [1] [Wikipedia: File system](http://en.wikipedia.org/wiki/File_system). URL: http://en.wikipedia.org/wiki/File_system (see p. 64).
- [2] [Wikipedia: HTree](http://en.wikipedia.org/wiki/HTree). URL: <http://en.wikipedia.org/wiki/HTree>.
- [3] D. Poirier. *Second Extended File System*. URL: <http://www.nongnu.org/ext2-doc> (see p. 62).
- [4] M. Gorman. “Chapter 11: Swap management”. In: *Understanding the Linux Virtual Memory Manager*. Prentice Hall, 2004. URL: <http://www.kernel.org/doc/gorman/> (see pp. 2–5).
- [5] R. Love. “Chapter 2: File I/O”. In: *Linux System Programming*. 2nd ed. O’Reilly, 2013 (see p. 64).
- [6] R. Love. “Chapter 8: File and directory management”. In: *Linux System Programming*. 2nd ed. O’Reilly, 2013 (see p. 64).
- [7] A. Silberschatz, P.B. Galvin, and G. Gagne. “Chapter 10: File system”. In: *Operating System Concepts*. 9th ed. Wiley, 2014 (see p. 64).
- [8] A. Silberschatz, P.B. Galvin, and G. Gagne. “Chapter 11: Implementing file systems”. In: *Operating System Concepts*. 9th ed. Wiley, 2014 (see p. 64).
- [9] A. Silberschatz, P.B. Galvin, and G. Gagne. “Chapter 12: Mass storage structure”. In: *Operating System Concepts*. 9th ed. Wiley, 2014 (see p. 64).
- [10] A.S. Tanenbaum and H. Bos. “Chapter 4: File systems”. In: *Modern Operating Systems*. 4th ed. Pearson, 2015 (see p. 64).
- [11] *Standard for Information Technology - Portable Operating System Interface (POSIX)*. Institute of Electrical and Electronics Engineers (IEEE) 1003.1-2008. 2008. URL: <http://standards.ieee.org> (see pp. 11–13).
- [12] M.K. McKusick et al. “A Fast File System for UNIX”. In: *ACM Transactions on Computer Systems (TOCS)* 2.3 (1984), pp. 181–197 (see p. 61).
- [13] V. Prabhakaran, A.C. Arpaci-Dusseau, and R.H. Arpaci-Dusseau. “Analysis and evolution of journaling file systems”. In: *USENIX Annual Technical Conference (ATEC)*. 2005, pp. 105–120.
- [14] K. Thompson. “UNIX Implementation”. In: *Bell System Technical Journal* 57.6 (1978), pp. 1931–1946 (see pp. 59, 60).
- [15] S.C. Tweedie. “Journaling the Linux ext2fs Filesystem”. In: *LinuxExpo*. 1998.